M1 Data Science, University of Lille
Statistics 2 - Lecture notes

Emilie Kaufmann (CNRS, Univ. Lille)
emilie.kaufmann@univ-lille.fr

March 26, 2024

Pre-requisite In statistics 1, you have seen:

- classical distributions

- examples of estimators

- confidence intervals

- the statistical testing protocol (type I and type II error)

- example of classical tests

In statistics 2, we will revisit statistical estimation and testing with a focus on *optimality*.
We will notably discuss:

- different performance measure for estimators

- generic estimation strategies, notably the maximum likelihood principle

- asymptotic properties of estimators

- likelihood-ratio based testing procedures

Several examples will come from an important family of distributions called exponential families.

Chapter 1

Estimation

1.1 Statistical inference

In statistical inference, we observe a realization of some random variable (or random vector) X, called the observation, whose distribution over some space $\mathcal{X}$ is P_X. The goal is to discover ("infer") some properties of this underlying distribution, assuming that P_X belongs to some set of possible distributions, called the *statistical model*. Depending on the situation, we may make assumptions on the cumulative distribution function (cdf) of X, F_X or on its density f_X with respect to some reference measure and the statistical model may be a set of distribution, a set of cdfs or a set of pdfs parameterized by some parameter θ:

$$\mathcal{M} = \{P_\theta, \theta \in \Theta\}, \quad \mathcal{M} = \{F_\theta, \theta \in \Theta\} \quad \text{or} \quad \mathcal{M} = \{f_\theta, \theta \in \Theta\}.$$

When the parameter space $\Theta \subseteq \mathbb{R}^d$, the model is called parametric, otherwise it is non-parametric. Given the "true" parameter θ (i.e. θ such that $P_X = P_\theta$), the probability space on which X is defined is denoted by $(\Omega, \mathcal{F}, \mathbb{P}_\theta)$, and the corresponding expectation is denoted by $\mathbb{E}_\theta$.

The n-sample example Often the random variable X is of the form $X = (X_1, \ldots, X_n)$ where the X_i are assumed to be iid realizations of the same distribution. These iid copies represent the repetition of some random experiment (for example the vote expressed by one individual in a population, or the effect of a treatment on one patient). These random variables X_i are defined on some common probability space $(\Omega, \mathcal{F}, \mathbb{P})$ and will most of the time take values in $\mathbb{R}$ (we could consider some multi-dimensional outcomes in, e.g. two-sample testing problems).

   In the n-sample setting, we denote by P the distribution of X_1 (which is the common distribution of all X_i's), by F the cdf of this distribution and by f its density (with respect to some reference measure ν), if it admits one. We will write indifferently

$$X_1, \ldots, X_n \overset{\text{i.i.d.}}{\sim} P, \quad X_1, \ldots, X_n \overset{\text{i.i.d.}}{\sim} F \quad \text{or} \quad X_1, \ldots, X_n \overset{\text{i.i.d.}}{\sim} f.$$

In that case, the statistical model is typically expressed as possible candidates for P, F or f. Those also denoted by P_θ, F_θ and f_θ, respectively (by a slight abuse of notation), for some parameter θ belonging to the parameter space Θ.

3

Example 1.1. *Take a Gaussian n-sample with known variance 1 and unknown mean $\theta \in \mathbb{R}$: $X_1, \ldots, X_n \overset{iid}{\sim} \mathcal{N}(\theta, 1)$. Let f_θ be the density of a $\mathcal{N}(\theta, 1)$ variable with respect to the Lebesgue measure (in $\mathbb{R}$):*

$$f_\theta(x) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{(x-\theta)^2}{2}\right).$$

If we look at the observation $X = (X_1, \ldots, X_n)$, the statistical model $\mathcal{M}$ is a set of multivariate Gaussian distributions whose densities with respect to the Lebesgue measure in $\mathbb{R}^n$ is

$$f_\theta(x_1, \ldots, x_n) = \prod_{i=1}^{n} f_\theta(x_i)$$

for some parameter $\theta \in \mathbb{R}$.

In statistical inference, we are interested in making statements about the "true" parameter θ generating the data or about some *parameter of interest* which can be some function of θ, denoted by $g(\theta)$. This statement can be a guess for its value (estimation), an interval to which it belongs (confidence interval) or the answer to some question about this parameter (statistical test).

Example 1.2. *$X_1, \ldots, X_n \overset{iid}{\sim} \mathcal{N}(\mu, \sigma^2)$. The parameter of the model is $\theta = (\mu, \sigma)$ and the parameter space is $\Theta = \{(\mu, \sigma), \mu \in \mathbb{R}, \sigma > 0\}$. If we are solely interested in estimating the mean, the parameter of interest is μ and σ may be called a nuisance parameter.*

In some situations, we may be interested in estimating more complex functions of θ. For example, assume that X_i models the amount of antibodies produced 15 days after receiving a vaccine. For a given disease, the vaccine is considered efficient if this amount exceeds some threshold v. A possible parameter of interest is the probability of efficacy of the vaccine, p, which can be expressed as

$$p = \mathbb{P}(X_1 \geq v) = 1 - \mathbb{P}(X_1 < v) = 1 - \mathbb{P}\left(\frac{X_1 - \mu}{\sigma} < \frac{v - \mu}{\sigma}\right) = 1 - \Phi\left(\frac{v - \mu}{\sigma}\right)$$

where Φ is the cdf of a $\mathcal{N}(0, 1)$ random variable.

Example 1.3 (regression model). *$Z_1, \ldots, Z_n \overset{iid}{\sim} P$. $X_i = (Z_i, Y_i) \in \mathcal{X} \times \mathbb{R}$ such that*

$$Y_i = h(Z_i) + \varepsilon_i$$

where $\varepsilon_i \overset{iid}{\sim} \mathcal{N}(0, 1)$ and $h : \mathcal{X} \to \mathcal{Y}$ is the regression function. The observation is $X = (X_1, \ldots, X_n)$ and the parameters of the model are P (that could belong to some parametric class of probability distributions) and the regression function h (that could belong to a parametric families of functions, e.g. linear functions). In that case the "parameter" of interest is usually the regression function.

1.2   Performance of an estimator

An estimator of $g(\theta)$ is any function of the observation $\widehat{g} = h(X)$ that is supposed to be "close" to the parameter of interest $g(\theta)$. When $X = (X_1, \ldots, X_n)$ has the n-sample structure, we will materialize the dependency in n of the estimator by writing $\widehat{g}_n = h(X_1, \ldots, X_n)$.

From its definition, $\widehat{g}$ is a random variable (or a random vector, when we estimate a multi-dimensional parameter), hence its quality will be expressed in terms of some properties of its distribution, which should ideally be concentrated around $g(\theta)$. Two important characteristics of this distributions are its mean and its variance.

Definition 1.4. *The* bias *of estimator $\widehat{g}$ of $g(\theta)$ is defined as* $\mathrm{b}_\theta(\widehat{g}) = \mathbb{E}_\theta[\widehat{g}] - g(\theta)$.
When $\mathrm{b}_\theta(\widehat{g}) = 0$, *the estimator is called* unbiased.

Definition 1.5. *The* variance *of a real-valued estimator $\widehat{g}$ is* $\mathrm{Var}_\theta[\widehat{g}] := \mathbb{E}_\theta[(\widehat{g} - \mathbb{E}_\theta[\widehat{g}])^2]$.

A good (real-valued) estimator has ideally a small bias and a small variance, which indicates that on average, its value is close to $g(\theta)$ and that under different realizations of the experiments, its value would not change too much. The closeness from $\widehat{g}$ to $g(\theta)$ can also directly be measured using their average distance, a notion that can also be meaningful in the multi-dimensional setting.

Definition 1.6. *The* quadratic risk *of an estimator $\widehat{g}$ of $g(\theta) \in \mathbb{R}^p$ is*

$$\mathrm{R}_\theta(\widehat{g}) = \mathbb{E}_\theta\left[\|\widehat{g} - g(\theta)\|^2\right],$$

where $\|u\|$ is the Euclidian norm in $\mathbb{R}^p$, such that $\|u\|^2 = u^\top u$. In the one-dimensional case ($p = 1$), this quantity is sometimes called the mean-squared error, *and denoted by* $\mathrm{MSE}_\theta(\widehat{g})$.

Theorem 1.7 (bias-variance decomposition). *Assume $g(\theta) \in \mathbb{R}$. We have*

$$\mathrm{R}_\theta(\widehat{g}) = (\mathrm{b}(\widehat{g}))^2 + \mathrm{Var}_\theta[\widehat{g}] \ .$$

Exercise 1.8. *Prove it.*

Comparing estimators with the quadratic risk   The quadratic risk can be used to compare estimators, and we say that an estimator $\widehat{g}$ is better than an estimator $\widetilde{g}$ if for all $\theta \in \Theta$, $\mathrm{R}_\theta(\widehat{g}) \leq \mathrm{R}_\theta(\widetilde{g})$. However, this relationship is not a total order, as there may exists estimators for which $\mathrm{R}_{\theta_1}(\widehat{g}) \leq \mathrm{R}_{\theta_1}(\widetilde{g})$ for some parameter θ_1 but $\mathrm{R}_{\theta_2}(\widehat{g}) > \mathrm{R}_{\theta_2}(\widetilde{g})$ for a different parameter θ_2.

Definition 1.9. *An estimator $\widehat{g}$ of $g(\theta)$ is called* admissible *is there exists no estimator $\widetilde{g}$ which is strictly better than $\widehat{g}_n$ i.e. for which*

$$\forall \theta \in \Theta, \ \ \mathrm{R}_\theta(\widetilde{g}) \leq \mathrm{R}_\theta(\widehat{g})$$

and the inequality is strict for at least one value θ_0.

Influence of the sample size   When X is a n-sample, the above properties for an estimator $\widehat{g}_n$ are all considering a fixed sample size n, and are not capturing another desirable property of an estimator: $\widehat{g}_n$ should get closer to $g(\theta)$ when the sample size n goes larger. We expect to $\widehat{g}_n$ to get closer to $g(\theta)$, meaning that its distribution concentrates for and more around $g(\theta)$. We will discuss these asymptotic properties in the next chapter.

Recap: Densities and Expectations   All the criteria for evaluating estimators in this section are expressed with expectations. In general, if Z is a random variable taking values in $\mathcal{Z}$ whose distribution P has a density f with respect to some reference measure ν, we have, for all function ϕ,

$$\mathbb{E}[\phi(Z)] = \int_\mathcal{Z} \phi(x) f(x) d\nu(x).$$

We will mostly see examples of random variables defined on $\mathcal{Z} = \mathbb{R}^d$ whose distributions have a density with respect to the Lebesgue measure in $\mathbb{R}^d$, or of discrete random random variables (i.e. for which $\mathcal{Z}$

is discrete) that have a density with respect to the counting measure. In the discrete case, the density is simply defined, for all $z \in \mathcal{Z}$, by

$$f(z) = P(\{z\}) = \mathbb{P}_{Z \sim P} (Z = z) \ .$$

Back to our statistical model, in the most common n-sample case in which $X = (X_1, \ldots, X_n) \overset{iid}{\sim} P_\theta$, we will often uncounter two cases.

Either $X_i \in \mathbb{R}$ and P_θ has a density with respect to the Lebesgue measure. Then

- for any $\phi : \mathbb{R}^n \to \mathbb{R}$, $\mathbb{E}_\theta[\phi(X)] = \int_\mathbb{R} \phi(x_1, \ldots, x_n) f_\theta(x_1, \ldots, x_n) dx_1 \ldots dx_n$

- for any $\phi : \mathbb{R} \to \mathbb{R}$, $\mathbb{E}_\theta[\phi(X_1)] = \int_\mathbb{R} \phi(u) f_\theta(u) du$

Or $X_i \in \mathcal{S}$ for some discrete set $\mathcal{S}$ (typically a subset of $\mathbb{N}$) and we have

- for any $\phi : \mathcal{S}^n \to \mathbb{R}$, $\mathbb{E}_\theta[\phi(X)] = \sum_{x \in \mathcal{S}^n} \phi(x_1, \ldots, x_n) f_\theta(x_1, \ldots, x_n)$

- for any $\phi : \mathcal{S} \to \mathbb{R}$, $\mathbb{E}_\theta[\phi(X_1)] = \sum_{u \in \mathcal{S}} \phi(u) f_\theta(u)$

1.3   Estimation procedures

1.3.1   The moment method

When $X = (X_1, \ldots, X_n) \overset{iid}{\sim} P_\theta$, the moment method can be used when the parameter of interest $g(\theta)$ can be expressed as a function of the moments of X_1.

In the simple case, we have

$$g(\theta) = \mathbb{E}_\theta\left[\phi(X_1)\right]$$

for some function ϕ such that $\mathbb{E}[|\phi(X_1)|] < \infty$. Motivated by the law of large numbers, we define the moment estimator

$$\widehat{g}_n := \frac{1}{n} \sum_{i=1}^{n} \phi(X_i)$$

which satisfies $\widehat{g}_n \to g(\theta)$, $\mathbb{P}_\theta$ – a.s.. Hence, this estimator is naturally going to be close to $g(\theta)$ at least for a large sample size n.

More generally, suppose that we seek to estimate a multi-dimensional parameter $\theta = (\theta_1, \ldots, \theta_k)$ and that for $1 \le j \le k$ the j-th moment can be expressed as some function of the parameter θ:

$$\mathbb{E}_\theta[X^j] = \alpha_j(\theta).$$

Letting $\widehat{\alpha}_j = \frac{1}{n} \sum_{i=1}^n X_i^j$ the j-th sample moment, the moment estimator is defined as the solution $\widehat{\theta}_n$ of the system of equations

$$\alpha_1(\theta) = \widehat{\alpha}_1, \ \ldots \ , \alpha_k(\theta) = \widehat{\alpha}_k.$$

Example 1.10. $X_1, \ldots, X_n \overset{iid}{\sim} \mathcal{N}(\mu, \sigma^2)$. *We can find the moment estimator for the parameter* $\theta = (\mu, \sigma^2)$. *There are two parameters so we can look at the first two moments.*

$$\begin{aligned} \mathbb{E}_\theta[X_1] &= \mu \\ \mathbb{E}_\theta[X_1^2] &= \mathrm{Var}_\theta[X_1] + (\mathbb{E}_\theta[X_1])^2 = \sigma^2 + \mu^2 \end{aligned}$$

The empirical first and second moments are $\frac{1}{n}\sum_{i=1}^{n} X_i$ and $\frac{1}{n}\sum_{i=1}^{n} X_i^2$ so we get the system

$$\begin{cases} \mu & = & \frac{1}{n}\sum_{i=1}^{n} X_i \\ \mu^2 + \sigma^2 & = & \frac{1}{n}\sum_{i=1}^{n} X_i^2 \end{cases}$$

from which we get $\widehat{\theta}_n = (\widehat{\mu}_n, \widehat{\sigma}_n^2)$ where

$$\widehat{\mu}_n = \frac{1}{n}\sum_{i=1}^{n} X_i \quad \text{and} \quad \widehat{\sigma}_n^2 = \frac{1}{n}\sum_{i=1}^{n} X_i^2 - \left(\frac{1}{n}\sum_{i=1}^{n} X_i\right)^2$$

We recognize the well-known empirical mean and empirical variance, which can also be rewritten

$$\widehat{\sigma}_n^2 = \frac{1}{n}\sum_{i=1}^{n} \left(X_i - \frac{1}{n}\sum_{k=1}^{n} X_k\right)^2.$$

1.3.2   The plug-in method

The plug-in method is also suited for the n-sample setting, when the parameter of interest can be expressed as some functional of P, the distribution of X_1 (for example some moment of this distribution, or some quantile), we write

$$g(\theta) = H(P).$$

The idea is construct some empirical version of this distribution, denoted by $\widehat{P}_n$ and to "plug-in" this empirical distribution, that is to define

$$\widehat{g}_n = H(\widehat{P}_n).$$

We now describe this empirical distribution.

Definition 1.11. *Given a n-sample $X = (X_1, \ldots, X_n) \in \mathbb{R}^n$, the empirical distribution $\widehat{P}_n$ is a probability measure on $\mathbb{R}$ defined as*

$$\widehat{P}_n = \frac{1}{n}\sum_{i=1}^{n} \delta_{X_i}$$

where δ_x is the Dirac measure in x. The Dirac measure is defined, for all event A as $\delta_x(A) = 1$ if $x \in A$, $\delta_x(A) = 0$ otherwise. For any $x \in \mathbb{R}$, we have

$$\widehat{P}_n(\{x\}) = \frac{1}{n}\sum_{i=1}^{n} \delta_{X_i}(\{x\}) = \frac{\#\{i : X_i = x\}}{n}.$$

$\widehat{P}_n$ *is a discrete distribution whose (finite) support (= set of values that have non-zero probability in the discrete case) is the values in $\{X_1, \ldots, X_n\}$.*

For any function ϕ, the expectation of $\phi(Z)$ when Z is distributed according to the empirical distribution $\widehat{P}_n$ is given by

$$\mathbb{E}_{Z \sim \widehat{P}_n}[\phi(Z)] = \sum_{x \in S} \phi(x)\widehat{P}_n(x) = \frac{1}{n}\sum_{i=1}^{n} \phi(X_i)$$

where S is the support of $\widehat{P}_n$, i.e. the number of distinct values in the observation X. In particular, the cdf $\widehat{F}_n$ of the empirical distribution, which by definition is $\widehat{F}_n(x) = \mathbb{P}_{Z \sim \widehat{P}_n}(Z \leq x) = \mathbb{E}_{Z \sim \widehat{P}_n}[\mathbb{1}(X_i \leq x)]$ can be written

$$\widehat{F}_n(x) = \frac{1}{n}\sum_{i=1}^{n} \mathbb{1}(X_i \leq x).$$

Remark 1.12. *When the functional $H(P)$ is defined as some expectation under P, the moment method and the plug-in method actually concide. Indeed, if*

$$g(\theta) = \mathbb{E}_{X \sim P}[\phi(X)]$$

the plug-in method yields

$$\widehat{g}_n = \mathbb{E}_{X \sim \widehat{P}_n}[\phi(X)] = \frac{1}{n} \sum_{i=1}^{n} \phi(X_i).$$

We also call such estimators "empirical estimators".

   But plug-in estimator can be more general when H is not defined as some expectation, for example we can define the empirical quantiles of a distribution to estimate its quantiles.

Exercise 1.13. *Using the plug-in approach, justify (again) the expression of the empirical mean and empirical variance of a distribution.*

1.3.3   Maximum Likelihood Estimation (MLE)

The maximum likelihood approach can be used to estimate $g(\theta) = \theta$ when the statistical model is of the form

$$\mathcal{M} = \{P_\theta : P_\theta \text{ has a density } f_\theta \text{ with respect to } \nu, \theta \in \Theta\}$$

where ν is a fixed reference measure (which is the same for all the distributions in the model). Such a model is called *dominated* (by the reference measure ν).

   In most practical cases, this reference measure will be the Lebesgue measure in $\mathbb{R}^d$ (when the distributions are continuous) or the counting measure on discrete set (when the distributions are discrete). In that case, the density is given by $f_\theta(x) = \mathbb{P}_\theta(X = x)$.

Definition 1.14. *The likelihood of the observation X given a parameter θ is defined by*

$$L(X; \theta) = f_\theta(X).$$

In the n-sample case, due to independence, the log-likelihood can be written

$$L(X_1, \ldots, X_n; \theta) = \prod_{i=1}^{n} f_\theta(X_i). \tag{1.1}$$

Example 1.15. *If $X_1, \ldots, X_n \sim \mathcal{B}(\theta)$. The density of a Bernoulli distribution with parameter θ can be written*

$$f_\theta(x) = \theta \mathbb{1}(x = 1) + (1 - \theta)\mathbb{1}(x = 0) = \theta^x (1 - \theta)^{1-x} \mathbb{1}(x \in \{0, 1\})$$

hence we have

$$L(X_1, \ldots, X_n; \theta) = \prod_{i=1}^{n} \theta^{X_i} (1 - \theta)^{1-X_i} = \theta^{\sum_{i=1}^{n} X_i} (1 - \theta)^{n - \sum_{i=1}^{n} X_i}$$

   If $X_1, \ldots, X_n \sim \mathcal{N}(\theta, \sigma^2)$, we get

$$L(X_1, \ldots, X_n; \theta) = \prod_{i=1}^{n} \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(X_i - \theta)^2}{2\sigma^2}\right)$$

which can also be re-arranged.

The likelihood can be interpreted as the probability of making observation X if the underlying parameter is θ. Indeed, if $\mathcal{M}$ is a set of discrete distributions (i.e. when ν is the counting measure), we have $f_\theta(x_i) = \mathbb{P}_\theta(X_i = x_i)$. Due to independence, we have

$$\mathbb{P}_\theta(X_1 = x_1, \ldots, X_n = x_n) = \prod_{i=1}^{n} f_\theta(x_i) = L(x; \theta) \quad \text{where } x = (x_1, \ldots, x_n)$$

In the continuous case (i.e. when ν is the Lebesgue measure), the probability of a given $x = (x_1, \ldots, x_n)$ is zero and we replace it by the value of the (joint) density in the point.

This observation motivates the maximum likelihood estimator as the estimator of $g(\theta) = \theta$ seeking the parameter θ for which the actual observation X is the most likely (i.e. which as it the largest "probability").

Definition 1.16. *A maximum likelihood estimator (MLE) of a parameter θ is an estimator satisfying*

$$\widehat{\theta} \in \underset{\theta \in \Theta}{argmax}\, L(X; \theta).$$

Remark 1.17. *As will be seen in some exercises, the maximum likelihood is not always unique.*

Computational considerations   In simple case, the maximum likelihood can be computed explicitely, by finding the critical points (for which the derivative are zero) and proving that it is indeed a maximizer (e.g., by checking that the second derivative is negative in the critical point). In more complex cases, it can only be approximated using some optimization algorithm. In complex models (like the Gaussian mixture model), more fancy approximation schemes are needed, like the EM algorithm (Expectation Maximization) algorithm.

From a computational prespective (and due to the common product form of the likelihood, see (1.1)) it is often more convenient to maximize the logarithm of the likelihood (which then becomes a sum).

Definition 1.18. *The log-likelihood of the observation X given a parameter θ is denoted by*

$$\ell(X; \theta) = \log L(X; \theta).$$

Exercise 1.19. *Poisson distributions are often used to model count data (e.g. the number of monthly purchases of a customer on an e-commerce website may follow a Poisson distribution). A Poisson distribution with parameter $\lambda > 0$, denoted by $\mathcal{P}(\lambda)$, is a discrete distribution defined as*

$$\mathbb{P}(Z = k) = \frac{\lambda^k}{k!} e^{-\lambda} \quad \text{for all } k \in \mathbb{N}.$$

Compute the maximum likelihood estimator of λ given iid observations $X_1, \ldots, X_n \overset{iid}{\sim} \mathcal{P}(\lambda)$. What other method(s) could you use to obtain the same estimator?

Example 1.20. *In the logistic regression model, there are iid pairs of observations (X_i, Y_i) where X_i comes from some distribution on $\mathbb{R}^d$ that is assumed to have some density and $Y_i \in \{-1, 1\}$ is such that*

$$\mathbb{P}(Y_i = 1 | X_i = x) = \frac{1}{1 + e^{-x^\top \theta}}$$

where $\theta \in \mathbb{R}^d$ is a regression parameter.

To define the likelihood of the data, we admit that the density of $(X_1, Y_1) \in \mathbb{R}^d \times \{0, 1\}$ is

$$f_\theta(x, y) = \mathbb{P}(Y_1 = y | X_1 = x) f(x).$$

You can verify that for all $x \in \mathbb{R}^d$ and all $y \in \{-1, 1\}$, $\mathbb{P}(Y_1 = y | X_1 = x) = \frac{1}{1 + e^{-yx^\top \theta}}$. The likelihood can therefore be written

$$L((X_1, Y_1), \ldots, (X_n, Y_n)) = \prod_{i=1}^n f(X_i) \left(\frac{1}{1 + e^{-Y_i(X_i^\top \theta)}} \right)$$

and a maximum likelihood estimator $\widehat{\theta}_n$ satisfies

$$\widehat{\theta}_n \in \underset{\theta \in \mathbb{R}^d}{argmin} \sum_{i=1}^n \log \left(1 + e^{-Y_i(X_i^\top \theta)} \right).$$

In this example, no closed-form expression exists for the MLE (unlike in a linear regression example), and we should resort to an optimization algorithm.

M-**estimators**   The MLE estimator is actually an example of a more general family of estimators called M-estimators, that are obtained as the minimization of some cumulative loss function of the data. A M estimator is of the form

$$\widehat{\theta}_n \in \underset{\theta \in \mathbb{R}^d}{argmin} \, M_n(\theta) \quad \text{where} \quad M_n(\theta) = \sum_{i=1}^n m(X_i; \theta).$$

In the particular case of the MLE, we have $m(X; \theta) = -\log f_\theta(X)$.

1.4   Beyond the likelihood

Under some additional regularity conditions on some dominated model it is possible to define an important quantity called the Fisher information, which is useful to provide a lower bound on the quality of an (unbiased) estimator (see Section 1.5). The Fisher information will also be useful in the next chapter to characterize the asymptotic distribution of the maximum likelihood estimator.

To ease the presentation, we define everything in the single-parameter setting, that is when the parameter space Θ is a subset of $\mathbb{R}$. All this concepts can be extended to the multi-dimensional setting by replacing derivative with gradients, variances with covariances, and second derivative with Hessian. We will briefly discuss this extension afterwards.

Definition 1.21. *A (uni-dimensional) parameteric model $\mathcal{M} = \{P_\theta, \theta \in \Theta \subseteq \mathbb{R}\}$ is regular if*

1. *it is dominated by some reference measure ν and for all θ, the support of f_θ, $S = \{x \in \mathcal{X} : f_\theta(x) > 0\}$ is independent of θ*

2. *for all $x \in S$, $\theta \mapsto f_\theta(x)$ is twice differentiable on Θ and its second derivative is continuous*

3. *for any event $\mathcal{E}$, we have*

$$\frac{\partial}{\partial \theta} \int_{\mathcal{E}} f_\theta(x) d\nu(x) = \int_{\mathcal{E}} \frac{\partial}{\partial \theta} f_\theta(x) d\nu(x)$$

$$\frac{\partial^2}{\partial^2 \theta} \int_{\mathcal{E}} f_\theta(x) d\nu(x) = \int_{\mathcal{E}} \frac{\partial^2}{\partial^2 \theta} f_\theta(x) d\nu(x)$$

Example 1.22. *We can show that many classical parameteric model satisfy this assumption (e.g. Bernoulli models, Gaussian model, Poisson model). A counter-example that will be studied in an exercise is the family of uniform distributions on $[0, \theta]$ for $\theta \in \mathbb{R}^+$, which already violates assumption 1.*

1.4.1 The Fisher information

Definition 1.23. *The score function is defined as the derivative of the log-likelihood.*

$$s(X; \theta) = \frac{\partial \ell(X; \theta)}{\partial \theta} = \frac{1}{f_\theta(X)} \frac{\partial f_\theta(X)}{\partial \theta}$$

An important property of the score under a regular model is the following.

Lemma 1.24. *Under a regular model, for all $\theta \in \Theta$, $\mathbb{E}_\theta[s(X; \theta)] = 0$.*

Proof.

$$\mathbb{E}_\theta[s(X; \theta)] \quad = \quad \int \frac{\partial \ell(x; \theta)}{\partial \theta} f_\theta(x) d\nu(x) = \int \frac{\frac{\partial}{\partial \theta} f_\theta(x)}{f_\theta(x)} f_\theta(x) d\nu(x) = \int \frac{\partial}{\partial \theta} f_\theta(x) d\nu(x)$$

$$\underset{(a)}{=} \int_S \frac{\partial}{\partial \theta} f_\theta(x) d\nu(x) \underset{(b)}{=} \frac{\partial}{\partial \theta} \left(\int_S f_\theta(x) d\nu(x) \right) \underset{(c)}{=} \frac{\partial}{\partial \theta} (1) = 0$$

where (a) uses property 1. of a regular model, (b) uses property 3 and (c) uses that f_θ is a density.

$\square$

The Fisher information matrix is defined as the variance of the score, which is equal to its second moment as the score is centered.

Definition 1.25. *In a regular model, the Fisher information of the observation X is defined as*

$$I^X(\theta) = \mathrm{Var}_\theta[s(X; \theta)] = \mathbb{E}_\theta\left[(s(X, \theta))^2\right].$$

In the n-sample case, we will write $I_n(\theta)$ to denote the Fisher information of the n-sample, and $I(\theta)$ the Fisher information of the observation made of a single realisation $X_1 \sim P_\theta$.

1.4.2 Some properties of the Fisher information

Lemma 1.26. *Under a regular model, it holds that $I^X(\theta) = -\mathbb{E}_\theta\left[\frac{\partial^2 \ell(X; \theta)}{\partial^2 \theta}\right]$.*

Exercise 1.27. *Prove it. Hint: start by computing the right-hand side, using property 3. of a regular model as in the proof of Lemma 1.24.*

The above lemma can be useful for the computation of the Fisher information. We now present another interesting property which is the additivity of the Fisher information. This property follows from the fact that the density of a couple of independent random variable is the product of their densities, and uses properties of the logarithm.

Lemma 1.28. *If X and Y are two independent random variables defined on the same probability space $(\Omega, \mathcal{F}, \mathbb{P}_\theta)$, then*

$$I^{(X,Y)}(\theta) = I^X(\theta) + I^Y(\theta) \ .$$

It follows that for a n sample $X = (X_1, \dots, X_n) \overset{iid}{\sim} P_\theta$,

$$I_n(\theta) = I^X(\theta) = nI^{X_1}(\theta) = nI(\theta) \ .$$

Example 1.29. *Consider the Bernoulli model $X_1, \dots, X_n \overset{iid}{\sim} \mathcal{B}(\theta)$. We have seen above that $I_n(\theta) = nI(\theta)$ where $I(\theta)$ is the Fisher information in a model with one Bernoulli observation X_1. In this model, we have*

$$
\begin{aligned}
L(X_1; \theta) &= \theta^{X_1}(1-\theta)^{1-X_1} \\
\ell(X_1; \theta) &= X_1 \log(\theta) + (1-X_1)\log(1-\theta) \\
\frac{\partial \ell(X_1; \theta)}{\partial \theta} &= \frac{X_1}{\theta} - \frac{1-X_1}{1-\theta} \\
\frac{\partial^2 \ell(X_1; \theta)}{\partial^2 \theta} &= -\frac{X_1}{\theta^2} + \frac{1-X_1}{(1-\theta)^2}
\end{aligned}
$$

hence

$$I(\theta) = -\mathbb{E}_\theta\left[\frac{\partial^2 \ell(X_1; \theta)}{\partial^2 \theta}\right] = \mathbb{E}_\theta\left[\frac{X_1}{\theta^2} - \frac{1-X_1}{(1-\theta)^2}\right] = \frac{1}{\theta} - \frac{1}{1-\theta} = \frac{1}{\theta(1-\theta)}$$

Finally, using Lemma 1.28, we get $I_n(\theta) = \frac{n}{\theta(1-\theta)}$.

Extension to the multi-dimensional setting   If $\theta = (\theta_1, \dots, \theta_d)$, the score is a vector in $\mathbb{R}^d$, defined as

$$s(X; \theta) = \nabla_\theta \ell(X; \theta) = \left(\frac{\partial \ell(X; \theta)}{\partial \theta_1}, \dots, \frac{\partial \ell(X; \theta)}{\partial \theta_d}\right)^\top \ .$$

In (an extension of the definition of a) regular model, the score satisfies $\mathbb{E}[s(X; \theta)] = 0$ and the Fisher information is defined as the (covariance) of the score, ie

$$I(\theta) = \mathbb{E}\left[(s(X, \theta))(s(X, \theta))^\top\right] \ .$$

The Fisher information is therefore a $d \times d$ matrix, and a counterpart of Lemma 1.26 can be proved:

$$I(\theta) = -\mathbb{E}\left[\left(\frac{\partial^2 \ell(X; \theta)}{\partial \theta_i \partial \theta_j}\right)_{\substack{1 \le i \le d \\ 1 \le j \le d}}\right] \ .$$

1.4.3   Interpretation of the Fisher information *(more advanced)*

The Fisher information will be shortly related to the minimal variance that a unbiased estimator can have. But we can still try to provide an explanation as to why it can be called "information".

First, due to its additivity property (Lemma 1.28), if we interpret $I(\theta)$ as an amount of "information" brought by one sample, we note that the Fisher information of a n-sample is the sum of all the information brought by individual samples. Moreover, another property is that given an observation X, any "summary" of this observation in the form of a statistic $S = s(X)$ has a smaller Fisher information.

Lemma 1.30. *For any statistic $S = s(X)$ of an observation X defined on a probability space $(\Omega, \mathcal{F}, \mathbb{P}_\theta)$, we have $I^S(\theta) \leq I^X(\theta)$.*

Proof. Let's write down the proof assuming that X takes values in a discrete space $\mathcal{X}$ (to avoid the concept of conditional density). X and $S = s(X)$ are clearly defined on the same probability space $(\Omega, \mathcal{F}, \mathbb{P}_\theta)$. We can write

$$\mathbb{P}_\theta(X = x) \quad = \quad \mathbb{P}_\theta(X = x, S = s(x)) = \mathbb{P}_\theta(X = x | S = s(x)) \mathbb{P}_\theta(S = s(x))$$

Hence, for any $x \in \mathcal{X}$, writing $s = s(x)$, we have

$$f_\theta(x) = f_\theta(x|s) \widetilde{f_\theta}(s)$$

where we introduce f_θ the density of X, $\widetilde{f_\theta}$ the density of S and $f_\theta(x|s) := \mathbb{P}_\theta(X = x | S = s)$. Taking the logarithm and differentiating twice yields

$$\frac{\partial^2 \log f_\theta(x)}{\partial^2 \theta} = \frac{\partial^2 \log \widetilde{f_\theta}(s)}{\partial^2 \theta} + \frac{\partial^2 \log f_\theta(x|s)}{\partial^2 \theta}$$

and in particular

$$\frac{\partial^2 \log f_\theta(X)}{\partial^2 \theta} = \frac{\partial^2 \log \widetilde{f_\theta}(S)}{\partial^2 \theta} + \frac{\partial^2 \log f_\theta(X|S)}{\partial^2 \theta}$$

Taking the expectation and using Lemma 1.26 yields

$$I^X(\theta) = I^S(\theta) - \mathbb{E}_\theta \left[\frac{\partial^2 \log f_\theta(X|S)}{\partial^2 \theta} \right]$$

We conclude by noting that

$$-\mathbb{E}_\theta \left[\frac{\partial^2 \log f_\theta(X|S)}{\partial^2 \theta} \right] = \sum_s \mathbb{P}_\theta(S = s) \underbrace{\left[-\mathbb{E}_\theta \left[\frac{\partial^2 \log \mathbb{P}_\theta(X|S = s)}{\partial^2 \theta} \right] \right]}_{\geq 0}$$

and the term between brackets is positive as it is the Fisher information of the conditionnal distribution of X given $(S = s)$.

$\square$

From this result a good statistic $S = s(X)$ is one that doesn't loose information, i.e. for which $I^S(\theta) = I^X(\theta)$. Sufficient statistic have this property, and are defined below.

Definition 1.31. *A statistic $S = s(X)$ is called sufficient for θ is the distribution of $X = (X_1, \ldots, X_n)$ contionally to S does not depend on θ.*

We admit the following characterization.

Theorem 1.32 (Neyman-Fisher). *The statistic $S = s(X_1, \ldots, X_n)$ is sufficient for θ is there exists two positive functions g and h such that the density of X can be written*

$$f_\theta(x_1, \ldots, x_n) = g(x_1, \ldots, x_n) h(s(x_1, \ldots, x_n); \theta) \,.$$

1.4.4   The Kullback-Leibler divergence

We define another information theoretic quantity that is related to the likelihood (or actually rather to a likelihood ratio) and provides some notion of "distance" (although it is not a distance in the topological sense) between probability measures.

Definition 1.33. *For two probability measure P and Q that have a densities f and g with respect to the same probability measure ν and such that $g(x) = 0 \Rightarrow f(x) = 0$, we have*

$$\mathrm{KL}(P, Q) = \mathbb{E}_{X \sim P}\left[\log \frac{f(x)}{g(x)}\right].$$

In particular, if P_θ and $P_{\theta'}$ are two distributions in a regular model (actually assumption 1. in Definition 1.21 is sufficient), we can define

$$\mathrm{K}(\theta, \theta') := \mathrm{KL}(P_\theta, P_{\theta'}) = \mathbb{E}_\theta\left[\log \frac{f_\theta(X)}{f_{\theta'}(X)}\right].$$

Example 1.34. *The KL divergence between $\mathcal{N}(\mu, \sigma^2)$ and $\mathcal{N}(\mu', \sigma^2)$ is*

$$\mathrm{K}(\mu, \mu') = \frac{(\mu - \mu')^2}{2\sigma^2}.$$

The KL divergence between two Bernoulli distributions of parameters θ and θ' is

$$\mathrm{K}(\theta, \theta') = \theta \log\left(\frac{\theta}{\theta'}\right) + (1 - \theta)\log\left(\frac{1 - \theta}{1 - \theta'}\right).$$

1.5   The Cramer-Rao lower bound

The Fisher information defined in the previous section enables us (in the case of uni-dimensional estimation) to solve the following question: what is the minimal variance of an unbiased estimator? We consider this question for regular models.

Theorem 1.35. *Assume the statistical model is regular. Let $\widehat{g}$ be an estimator of $g(\theta) \in \mathbb{R}$ where g is differentiable. We assume that $\widehat{g} = h(X)$ is such that $\mathbb{E}_\theta[\widehat{g}_n] = g(\theta)$ (unbiased estimator) and*

$$\frac{\partial}{\partial \theta}\int h(x)f_\theta(x)d\nu(x) = \int h(x)\left(\frac{\partial}{\partial \theta}f_\theta(x)\right)d\nu(x)$$

Then, for all $\theta \in \Theta$,

$$\mathrm{Var}_\theta[\widehat{g}] \geq \frac{(g'(\theta))^2}{I(\theta)}.$$

Proof. The idea of the proof is to differentiate $g(\theta) = \mathbb{E}_\theta[h(X)]$ and introduce the score. Using one of

the assumptions, we can write

$$
\begin{aligned}
g'(\theta) \;=\; & \frac{\partial}{\partial\theta}\int_S h(x)f_\theta(x)d\nu(x) = \int_S h(x)\left(\frac{\partial}{\partial\theta}f_\theta(x)\right)d\nu(x) \\
=\; & \int_S h(x)\left(\frac{\partial}{\partial\theta}\log f_\theta(x)\right)f_\theta(x)d\nu(x) \\
\stackrel{(a)}{=}\; & \int_S \left(h(x)-\mathbb{E}_\theta[h(X)]\right)\left(\frac{\partial}{\partial\theta}\log f_\theta(x)\right)f_\theta(x)d\nu(x) + \mathbb{E}_\theta[h(X)]\underbrace{\int_S\left(\frac{\partial}{\partial\theta}\log f_\theta(x)\right)f_\theta(x)d\nu(x)}_{=0} \\
\stackrel{(b)}{=}\; & \int_S \left(h(x)-\mathbb{E}_\theta[h(X)]\right)\left(\frac{\partial}{\partial\theta}\log f_\theta(x)-\mathbb{E}_\theta\left[\frac{\partial}{\partial\theta}\log f_\theta(X)\right]\right)f_\theta(x)d\nu(x)
\end{aligned}
$$

where both (a) and (b) use that the expected score is zero by Lemma 1.24.

Now we assume that $\mathbb{E}_\theta[h^2(X)] < \infty$ (otherwise, the inequality in Theorem 1.35 is trivially true). Then we can use the Cauchy-Schwarz inequality to get

$$
\begin{aligned}
|g'(\theta)| \;\leq\; & \sqrt{\mathbb{E}_\theta\left[(h(x)-\mathbb{E}_\theta[h(X)])^2\right]}\sqrt{\mathbb{E}_\theta\left[\left(\frac{\partial}{\partial\theta}\log f_\theta(x)-\mathbb{E}_\theta\left[\frac{\partial}{\partial\theta}\log f_\theta(X)\right]\right)^2\right]} \\
\leq\; & \sqrt{\mathrm{Var}_\theta[h(X)]}\sqrt{\mathrm{Var}_\theta\left[\frac{\partial}{\partial\theta}\log f_\theta(X)\right]} \\
\leq\; & \sqrt{\mathrm{Var}_\theta[h(X)]}\sqrt{I(\theta)}
\end{aligned}
$$

where the last step uses the definition of the Fisher information.

$\square$

An unbiased estimator that achieves the Cramer-Rao lower bound for all values of $\theta \in \Theta$ is called efficient (or uniformly efficient). The example below show that there exists efficient estimators.

Exercise 1.36. *Verify that in the Bernoulli model $X_1,\ldots X_n \stackrel{iid}{\sim} \mathcal{B}(p)$ the MLE is an efficient estimator.*

1.6 Exponential families

Actually, the reason why in the Bernoulli model we can find an efficient estimator comes from the fact that the set of Bernoulli distributions is a particular example of exponential family. We define exponential families below, and we will see several of their properties in this class.

Definition 1.37. *An exponential family is a set of probability distibutions on some set $\mathcal{X}$ defined as*

$$
\mathcal{P} = \left\{P_\theta, \theta \in \Theta : \; P_\theta \text{ has a density } \; f_\theta(x) = h(x)\exp\left(a(\theta)^\top T(x)-b(\theta)\right) \text{ wrt to } \nu\right\}
$$

where ν is a reference measure (common to all distributions), $h : \mathcal{X} \to \mathbb{R}^+$ is a positive function, $a : \Theta \to \mathbb{R}^d$, $b : \Theta \to \mathbb{R}$ and $T : \mathcal{X} \to \mathbb{R}^d$ are some functions and $u^\top v = \sum_{i=1}^d u_i v_i$ is the scalar product in $\mathbb{R}^d$.

$T(x) \in \mathbb{R}^d$ is called the canonical statistic and d is the dimension of the exponential family. In a one-dimensional exponential family, the density can simply be expressed

$$
f_\theta(x) = h(x)\exp\left(a(\theta)T(x)-b(\theta)\right).
$$

Exercise 1.38. *Justify that family of Bernoulli distribution $\mathcal{P} = \{\mathcal{B}(p), p \in (0,1)\}$ form an exponential family (of dimension 1).*

Actually, we can prove that efficient estimator can only exist in some exponential families, and for a particular parameter to estimate. There are therefore not so much common. In the next chapter, we will define an asymptotic notion of efficiency, which can be easier to attain.

Chapter 2

Asymptotic properties of estimators

In this chapter, we focus on the n-sample case, in which $X = (X_1, \ldots, X_n) \overset{iid}{\sim} P_\theta$. For each n, given an estimator $\widehat{g}_n = h(X_1, \ldots, X_n)$ of a certain parameter of interest $g(\theta)$, we are interested in studying the sequence of estimators $(\widehat{g}_n)_n$ when the sample size n grows large. As the $\widehat{g}_n$ are random variables, we first recap the different notion of convergences, as well as some important results.

2.1 Refresher: Convergence of random variables

Definition 2.1. *Let $Z_1, Z_2, \ldots$ be a sequence of random variable and let Z be another random variable. Let F_n denote the CDF of Z_n and let F denote the cdf of Z.*

1. **Z_n converges to Z in probability** *if, for every $\varepsilon > 0$, $\lim_{n \to \infty} \mathbb{P}(|Z_n - Z| > \varepsilon) = 0$.*

   We write $Z_n \overset{P}{\longrightarrow} Z$.

2. **Z_n converges to Z in distribution** *if, $\lim_{n \to \infty} F_n(t) = F(t)$ for all t for which F is continuous.*

   We write $Z_n \rightsquigarrow Z$.

3. **Z_n converges to Z almost surely** *if $\mathbb{P}\left(\lim_{n \to \infty} Z_n = Z\right) = 1$. We write $Z_n \overset{a.s.}{\longrightarrow} Z$.*

4. **Z_n converges to Z in quadratic mean** *if $\lim_{n \to \infty} \mathbb{E}\left[(Z_n - Z)^2\right] = 0$. We write $Z_n \overset{L^2}{\longrightarrow} Z$.*

In statistics, the first two notions are the most common, and we will mostly discuss them in the following. The definitions above were all given for real-values random variables, but can be extended to the multi-dimensional setting. For the convergence in probability, the distance between Z_n and Z and $\mathbb{R}^d$ can no longer be measured with the absolute value, but given any distance d on $\mathbb{R}^d$ (for example the Euclidian distance), we define $Z_n \overset{P}{\longrightarrow} Z$ is for all $\varepsilon > 0$, $\lim_{n \to \infty} \mathbb{P}\left(d(Z_n, Z) > \varepsilon\right) = 0$.

The convergence in distribution in $\mathbb{R}^d$ can still be characterized by the cdf, but in this case, the cdf is a multi-variate function and we should have, for all $z = (z_1, \ldots, z_d)$ in which F is continuous,

$$\lim_{n \to \infty} \mathbb{P}\left(Z_n^1 \le z_1, \ldots, Z_n^d \le z_d\right) = \mathbb{P}(Z^1 \le z_1, \ldots, Z^d \le z_d) = 0.$$

Example 2.2. *$Z_n \sim \mathcal{N}(0, \frac{1}{n})$. Justify that Z_n converges to 0 (the random variable that is constant and equal to zero) in distribution and in probability.*

2.1.1 Properties

The following relationship between the different convergence notions are useful.

Lemma 2.3.   *1.* $X_n \xrightarrow{P} X$ *implies that* $X_n \rightsquigarrow X$

   2. $X_n \xrightarrow{P} c$ *where c is a constant if and only if* $X_n \rightsquigarrow X$

   3. $X_n \xrightarrow{L^2} X$ *implies that* $X_n \xrightarrow{P} X$

We note that (a) and (c) are not equivalences. In particular, beyond the case of convergence to constants, the convergence in distribution does not imply the convergence in probability. A (contrived) counter-example is the following: take any symmetric distribution Y, that is a distribution for which Y and $-Y$ have the same distribution (for example, a centered Gaussian distribution). Define $Z_n = Y$ for all n and $Z = -Y$. As the cdf and Z_n and that of Z are equal, we have in particular $Z_n \rightsquigarrow Z$. However, $\mathbb{P}(|Z_n - Z| > \varepsilon) = \mathbb{P}(|2Y| > \varepsilon)$ does not converge to zero for every ε (unless $Y = 0$ a.s.).

Lemma 2.4 (continuous mapping). *Let* $g : \mathcal{X} \to \mathbb{R}$ *be a continuous function. Then*

   - *If* $X_n \xrightarrow{P} X$ *then* $g(X_n) \xrightarrow{P} g(X)$

   - *If* $X_n \rightsquigarrow X$ *then* $g(X_n) \rightsquigarrow g(X)$

Lemma 2.5 (Slutsky lemma). *If* $X_n \rightsquigarrow X$ *and* $Y_n \rightsquigarrow c$ *where c is a constant, then, for any continuous function g,*

$$g(X_n, Y_n) \rightsquigarrow g(X, c).$$

In particular

   - $X_n + Y_n \rightsquigarrow X + c$

   - $X_n Y_n \rightsquigarrow cX$

Slutsky's lemma is a consequence of the fact that as a couple of random variables (X_n, Y_n) converges in distribution to (X, c) (and the fact that the continuous mapping lemma also apply to multi-variate random variables).

2.1.2 Two fundamental theorems

We recall here the two fundamental theorems in statistics: the law of large numbers and the central limit theorem. Given an iid sequence Z_i, they provide some convergence results for the empirical average

$$\widehat{Z}_n = \frac{1}{n} \sum_{i=1}^n Z_i.$$

Theorem 2.6 (Law of large numbers). *If* $(Z_i)_{i \in \mathbb{N}}$ *is an iid sequence with* $\mathbb{E}[Z_1] < \infty$*, we have*

$$\frac{1}{n} \sum_{i=1}^n Z_i \xrightarrow{P} \mathbb{E}[Z_1]$$

Actually, a stronger version of this result (called the strong law of large numbers) holds under the same assumptions, in which the convergence in probability is replaced by an almost sure convergence.

Theorem 2.7 (Central limit theorem). *If $(Z_i)_{i \in \mathbb{N}}$ is an iid sequence with $\mathbb{E}[Z_1^2] < \infty$, letting $\mu = \mathbb{E}[Z_1]$ and $\sigma^2 = \mathrm{Var}[Z_1]$, we have*

$$\sqrt{\frac{n}{\sigma^2}} \left(\widehat{Z}_n - \mu \right) \rightsquigarrow \mathcal{N}(0,1)$$

Under the hypotheses of the central limit theorem, $\widehat{Z}_n$ can be written

$$Z_n = \mu + \sqrt{\frac{\sigma^2}{n}} Y_n$$

where $Y_n \rightsquigarrow \mathcal{N}(0,1)$. Therefore, informally, the distribution of $\widehat{Z}_n$ is close to $\mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right)$, a Gaussian distribution whose variance decays to zero and is therefore more and more concentrated around μ. We may write $\widehat{Z}_n \approx \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right)$ and talk about the asymptotic distribution of $\widehat{Z}_n$.

2.2 Consistency and asymptotic normality

Definition 2.8. *An estimator $\widehat{g}_n$ of $g(\theta)$ is consistent if for every $\theta \in \Theta$, $\widehat{g}_n \xrightarrow{P} g(\theta)$.*

Consistency of estimators will often follow from the law of large numbers. When we further have an almost sure convergence, that is when $\widehat{g}_n \xrightarrow{a.s.} g(\theta)$, we shall say that $\widehat{g}_n$ is strongly consistent.

Lemma 2.3 and Lemma 2.4 also yield the following properties:

- If the quadratic risk $\mathrm{R}_\theta(\widehat{g}_n)$ goes to zero when n goes to infinity, $\widehat{g}_n$ is consistent.

- If $\widehat{\theta}_n$ is a consistent estimator of θ and g is a continuous mapping, then $\widehat{g}_n = g(\widehat{\theta}_n)$ is a consistent estimator of $g(\theta)$.

Example 2.9. *Justify that the empirical mean and empirical variance defined in the previous chapter are consistent estimators.*

Given a consistent estimator $\widehat{g}_n$, we may be interested in how fast $\widehat{g}_n - g(\theta)$ converges to zero. To do so, we will look at the limit distribution of (some re-normalization) of this random variable.

If we take the example of the empirical $\widehat{\mu}_n = \frac{1}{n} \sum_{i=1}^{n} X_i$ estimator of the common mean μ of some n sample $(X_1, \ldots, X_n)$ that has variance σ^2, the Central Limit Theorem tells us that

$$\sqrt{n} \left(\widehat{\mu}_n - \mu \right) \rightsquigarrow \mathcal{N}(0, \sigma^2)$$

Here, the limit distribution is Gaussian, and the convergence speed is $\sqrt{n}$. Due to the generality of the Central Limit Theorem, we expect this Gaussian limit behavior to be a general pattern for estimators, hence the definition of asymptotic normality.

Definition 2.10. *An estimator is $\widehat{g}_n$ of $g(\theta)$ is asymptotically normal is it satisfies, for all $\theta \in \Theta$,*

$$\sqrt{n} \left(\widehat{g}_n - g(\theta) \right) \rightsquigarrow \mathcal{N}(0, \sigma_\theta^2)$$

where σ_θ^2 is called the asymptotic variance.

It is worth mentioning that there exists estimators that are have a limiting distribution, but are not asymptotically normal. It means that they satisfy something like

$$g(n)\left(\widehat{g}_n - g(\theta)\right) \rightsquigarrow Z$$

where $g(n)$ is some convergence speed (that can be different than $\sqrt{n}$) and Z is some fixed distribution (that is not necessarily Gaussian).

Example 2.11. *As studied in exercise, in the model $X_1, \ldots, X_n \overset{iid}{\sim} \mathcal{U}([0, \theta])$, the moment estimator is $\widehat{\theta}_n = \frac{2}{n}\sum_{i=1}^n X_i$ while the MLE is $\widetilde{\theta}_n = \max_{i=1..n} X_i$.*

Using the Central Limit Theorem (and the continuous mapping lemma), we can show that

$$\sqrt{n}\left(\widehat{\theta}_n - \theta\right) \rightsquigarrow \mathcal{N}\left(0, \frac{\theta^2}{3}\right)$$

hence the moment estimator is asymptotically normal with asymptotic variance $\sigma_\theta = \frac{\theta^2}{3}$.

On the other hand, we computed the distribution of $\widetilde{\theta}_n$ in exercise, showing that

$$\mathbb{P}\left(\widetilde{\theta}_n \le x\right) = \begin{cases} 1 & \text{if } x \ge \theta \\ \frac{x^n}{\theta^n}\mathbb{1}_{[0,\theta]}(x) & \text{else.} \end{cases}$$

Hence we have, for all $u > 0$,

$$\mathbb{P}\left(\theta - \widetilde{\theta}_n \ge u\right) = \left(1 - \frac{u}{\theta}\right)^n \mathbb{1}_{[0,\theta]}(u)$$

and finally, for all $t > 0$,

$$\mathbb{P}\left(n\left(\theta - \widetilde{\theta}_n\right) \ge t\right) = \left(1 - \frac{t}{n\theta}\right)^n \mathbb{1}_{[0,n\theta]}(t)$$

The limit of the right-hand side when n goes to infinity is equal to $e^{-\frac{t}{\theta}} = \mathbb{P}(Z > t)$ where Z is an exponential distribution with parameter $1/\theta$. Finally, one can write

$$n\left(\theta - \widetilde{\theta}_n\right) \rightsquigarrow \mathcal{E}\left(\theta^{-1}\right).$$

This provide another argument for using the MLE over the moment estimator in this particular case, as its asymptotic convergence is faster.

In Section 2.4, we will actually see that for regular models, the MLE is asymptotically Gaussian. The reason for this different behavior stems from the fact that the model considered here is not regular: one can indeed see that all the possible densities do not have the same support.

Comparing asymptotically normal estimators   Between two asymptotically normal estimator, the one with smallest asymptotic variance σ_θ^2 is the one that converges "faster" to the parameter $g(\theta)$. This can be measured by the fact that, if we build asymptotic confidence intervals for $g(\theta)$ of level $1 - \alpha$, using the estimator with smallest asymptotic variance will yield the smallest confidence region.

If two asymptotically normal estimators $\widehat{g}_n$ and $\widetilde{g}_n$ have respective asymptotic variances σ_θ^2 and $\widetilde{\sigma}_\theta^2$ and that $\sigma_\theta^2 \le \widetilde{\sigma}_\theta^2$ for all $\theta \in \Theta$ (with at least one strict inequality), we say that $\widehat{g}_n$ is asymptotically more efficient than $\widetilde{g}_n$.

Remark 2.12. *Asymptotic normality is a stronger notion than consistency. Indeed, any asymptotically normal estimator is also consistent by the following argument. Fix $\varepsilon > 0$. We have*

$$\mathbb{P}\left(|\widehat{g}_n - g(\theta)| > \varepsilon\right) = \mathbb{P}\left(\sqrt{n}|\widehat{g}_n - g(\theta)| > \varepsilon\sqrt{n}\right).$$

For all $\alpha > 0$, there exists n_0 such that n larger than n_0 implies $\varepsilon\sqrt{n} \geq \Phi^{-1}(1 - \alpha/4)$ where Φ is the standard normal cdf. For $n \geq n_0$, it therefore holds that

$$\mathbb{P}\left(|\widehat{g}_n - g(\theta)| > \varepsilon\right) \leq \mathbb{P}\left(\sqrt{n}|\widehat{g}_n - g(\theta)| > \Phi^{-1}(1 - \alpha/4)\right) \xrightarrow[n\to\infty]{} \mathbb{P}_{Z\sim\mathcal{N}(0,1)}\left(Z > \Phi^{-1}(1 - \alpha/4)\right) = \frac{\alpha}{2}.$$

Hence, for any $\alpha > 0$, for n large enough, $\mathbb{P}\left(|\widehat{g}_n - g(\theta)| > \varepsilon\right) \leq \alpha$. It follows that $\lim_{n\to\infty} \mathbb{P}\left(|\widehat{g}_n - g(\theta)| > \varepsilon\right) = 0$ and $\widehat{g}_n \xrightarrow{P} g(\theta)$.

2.3 The Delta method

We now present a useful tool to compute asymptotic distributions of some transformation of an asymptotically normal estimator: the so-called Delta method. This result implies that under some mild conditions, if $\widehat{\theta}_n$ is an asymptotically normal estimator of θ, then $g(\widehat{\theta}_n)$ is an asymptotically normal estimator of $g(\widehat{\theta}_n)$.

Theorem 2.13. *Suppose that for some sequence of random variance (Z_n),*

$$\sqrt{n}(Z_n - \mu) \rightsquigarrow \mathcal{N}(0, \sigma^2)$$

and that g is a differentiable function such that $g'(\mu) \neq 0$. Then

$$\sqrt{n}(g(Z_n) - g(\mu)) \rightsquigarrow \mathcal{N}\left(0, (g'(\mu))^2\sigma^2\right).$$

In other words,

$$Z_n \approx \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right) \quad \text{implies that} \quad g(Z_n) \approx \mathcal{N}\left(g(\mu), (g'(\mu))^2\frac{\sigma^2}{n}\right).$$

Proof. The proof follows from using a Taylor expansion around μ. As g is differentiable, we have that for all n, there exists μ_n in (Z_n, μ) (if $Z_n < \mu$) or in (μ, Z_n) (if $Z_n \geq \mu$) such that

$$g(Z_n) = g(\mu) + g'(\mu_n)(Z_n - \mu)$$

hence

$$\sqrt{n}|g(Z_n) - g(\mu)| = g'(\mu_n)\sqrt{n}(Z_n - \mu).$$

As $|\mu_n - \mu| \leq |Z_n - \mu|$ and $Z_n \xrightarrow{P} \mu$ (by the same argument used in Remark 2.12), we have that $\mu_n \xrightarrow{P} \mu$. If we assume g' to be continuous[1], it follows from Lemma 2.4 that $g'(\mu_n) \xrightarrow{P} g'(\mu)$.

By assumption, we also have that $\sqrt{n}(Z_n - \mu) \rightsquigarrow Z$ where $Z \sim \mathcal{N}(0, 1)$. It follows from Slutsky's lemma (Lemma 2.5) that

$$\sqrt{n}|g(Z_n) - g(\mu)| \rightsquigarrow g'(\mu)Z$$

whose distribution is $\mathcal{N}(0, (g'(\mu))^2\sigma^2)$.

---

[1] A slightly more complicated proof can also be given when g is not continuous, see e.g. [Rivoirard and Stoltz, 2009]

$\square$

There exists also a multi-variate version of the Delta method, stated below.

Theorem 2.14. *Let $Z_n = (Z_{n,1}, \ldots, Z_{n,d})$ be a sequence of random vectors in $\mathbb{R}^d$ such that*

$$\sqrt{n}(Z_n - \mu) \rightsquigarrow \mathcal{N}(0, \Sigma)$$

where $\mu \in \mathbb{R}^d$ and $\Sigma \in \mathbb{R}^{d \times d}$ is the covariance matrix. Let $g : \mathbb{R}^d \to \mathbb{R}$ be a differentiable function and let $\nabla g(z) = (\frac{\partial g}{\partial z_1}(z), \ldots, \frac{\partial g}{\partial z_d}(z))^\top$ be its gradient. If all the components of $\nabla g(\mu)$ are non-zero, then

$$\sqrt{n}\left(g(Z_n) - g(\mu)\right) \rightsquigarrow \mathcal{N}\left(0, (\nabla g(\mu))^\top \Sigma \left(\nabla g(\mu)\right)\right)$$

2.4   Asymptotic properties of the Maximum Likelihood Estimator

Given an iid sample $X_1, \ldots, X_n \sim P_\theta$, we recall that the maximum likelihood estimator of the parameter θ is defined

$$\widehat{\theta}_n \in \underset{\theta \in \Theta}{\operatorname{argmax}}\ L(X_1, \ldots, X_n; \theta).$$

Despite its implicit definition, as the maximizer of some function, we will see that this estimator enjoys strong asymptotic performance guarantees, when the model satisfies some assumptions. In particular, we will assume that the model is identifiable, that is

$$\forall (\theta, \theta') \in \Theta, P_\theta = P_{\theta'} \quad \text{if and only if} \quad \theta = \theta'.$$

2.4.1   Rationale

First, let us try to understand why the MLE is a good estimator. Let us denote by $\theta_\star$ the true parameter from which the data is generated. The maximum likelihood can be rewritten as follows, introducing artificially the likelihood under $\theta_\star$. Indeed, one can write

$$
\begin{aligned}
\widehat{\theta}_n &\in \underset{\theta \in \Theta}{\operatorname{argmax}}\ \frac{f_\theta(X_1) \ldots f_\theta(X_n)}{f_{\theta_\star}(X_1) \ldots f_{\theta_\star}(X_n)} \\
\widehat{\theta}_n &\in \underset{\theta \in \Theta}{\operatorname{argmin}}\ \frac{f_{\theta_\star}(X_1) \ldots f_{\theta_\star}(X_n)}{f_\theta(X_1) \ldots f_\theta(X_n)} \\
\widehat{\theta}_n &\in \underset{\theta \in \Theta}{\operatorname{argmin}}\ \log \frac{f_{\theta_\star}(X_1) \ldots f_{\theta_\star}(X_n)}{f_\theta(X_1) \ldots f_\theta(X_n)} \\
\widehat{\theta}_n &\in \underset{\theta \in \Theta}{\operatorname{argmin}}\ \sum_{i=1}^n \log\left(\frac{f_{\theta_\star}(X_i)}{f_\theta(X_i)}\right) \\
\widehat{\theta}_n &\in \underset{\theta \in \Theta}{\operatorname{argmin}}\ \frac{1}{n}\sum_{i=1}^n \log\left(\frac{f_{\theta_\star}(X_i)}{f_\theta(X_i)}\right)
\end{aligned}
$$

Hence, $\widehat{\theta}_n$ can be rewritten as the minimizer of some empirical average. Introducing the notation

$$M_n(\theta, \theta_\star) = \frac{1}{n}\sum_{i=1}^n \log\left(\frac{f_{\theta_\star}(X_i)}{f_\theta(X_i)}\right)$$

we know by the law of large number that, for all $\theta \in \Theta$,

$$M_n(\theta, \theta_\star) \xrightarrow{P} \mathbb{E}_{\theta_\star}\left[\log\left(\frac{f_{\theta_\star}(X_1)}{f_\theta(X_1)}\right)\right] = \mathrm{KL}\left(P_{\theta_\star}, P_\theta\right) \tag{2.1}$$

where $\mathrm{KL}(P, P')$ is the KL divergence between distributions, introduced in Definition 1.33. The KL divergence is not a distance but it still satisfies the following important property: $\mathrm{KL}(P, P') = 0$ if an only if $P = P'$. In particular, $\mathrm{KL}(P_{\theta_\star}, P_\theta) = 0$ if and only if $P_{\theta_\star} = P_\theta$, i.e. $\theta = \theta_\star$ as the model is identifiable. Thus we have

$$\underset{\theta \in \Theta}{\mathrm{argmin}}\, \mathrm{KL}\left(P_{\theta_\star}, P_\theta\right) = \theta_\star.$$

Hence, our hope is to prove that, under the model $\mathbb{P}_{\theta_\star}$,

$$\widehat{\theta}_n = \underset{\theta \in \Theta}{\mathrm{argmin}}\, M_n(\theta, \theta_\star) \xrightarrow{P} \underset{\theta \in \Theta}{\mathrm{argmin}}\, \mathrm{KL}\left(P_{\theta_\star}, P_\theta\right) = \theta_\star$$

This will require slightly more sophisticated arguments than the convergence of the objective function to minimize given in (2.1). We present them in the next section for more general M-estimators, that are also expressed as minimizer of empirical averages.

2.4.2 Consistency of M-estimators

A M-estimator is any estimator defined as a minimizer of some empirical average:

$$\widehat{\theta}_n = \underset{\theta \in \Theta}{\mathrm{argmin}}\, M_n(\theta) \quad \text{with} \quad M_n(\theta) = \frac{1}{n}\sum_{i=1}^n m_\theta(X_i)$$

Letting $M(\theta) = \mathbb{E}[m_\theta(X_1)]$, if this expectation is finite we have $M_n(\theta) \xrightarrow{P} M(\theta)$ for all $\theta \in \Theta$, and we hope that $\widehat{\theta}_n$, a minimizer of $M_n(\theta)$, converges to $\theta_0 = \underset{\theta \in \Theta}{\mathrm{argmin}}\, M(\theta)$.

Example 2.15. *In supervised learning, we observe iid pairs of the form (X_i, Y_i) coming from some unknown distribution P on $\mathcal{X} \times \mathcal{Y}$ where $\mathcal{X}$ is the feature space, often $\mathbb{R}^d$ and label space which is either finite (classification) or continuous (regression). The goal is to produce a predictor $\widehat{f}_n : \mathcal{X} \to \mathcal{Y}$, which is a data-dependent function mapping the feature to the label. Due to the generic empirical risk minimization principle, many predictor can be expressed as M-estimators.*

Given a class of function $\mathcal{F}$, and some loss function $L : \mathcal{Y} \times \mathcal{Y} \to \mathbb{R}_+$, we can define

$$\widehat{f}_n \in \underset{f \in \mathcal{F}}{\mathrm{argmin}}\, \frac{1}{n}\sum_{i=1}^n L(f(X_i), Y_i).$$

In this general (non-parametric) setting, the "parameter" is a function f (possible predictor), and we have $m_f((X_i, Y_i)) = L(f(X_i), Y_i)$. We hope that $\widehat{f}_n$ converges to $f_0 \in \underset{f \in \mathcal{F}}{\mathrm{argmin}}\, M(f)$ where $M(f) = \mathbb{E}_{(X,Y) \sim P}[L(f(X), Y)]$, that is to a predictor that minimizes the risk associated to the loss function L.

Sometimes, the class of function $\mathcal{F}$ can be described by a small set of parameters (e.g. a set of linear functions) and the regressor obtained by ridge regression can be defined as $\widehat{f}_n(x) = \widehat{\theta}_n^\top x$ for $x \in \mathbb{R}^d$ where

$$\widehat{\theta}_n = \underset{\theta \in \mathbb{R}^d : \|\theta\|_2 \leq C}{\mathrm{argmin}}\, \frac{1}{n}\sum_{i=1}^n (Y_i - \theta^\top X_i)^2\,.$$

In this case we hope that $\widehat{\theta}_n$ is close to $\theta_0 \in \underset{\theta \in \mathbb{R}^d, \|\theta\| \leq C}{\mathrm{argmin}}\, \mathbb{E}_{(X,Y) \sim P}[(Y - \theta^\top X)^2]$.

To establish the convergence of $\widehat{\theta}_n$ to θ_0, we need two properties. The first one is a property of the minimizer θ_0, which has to be a strict local minima, and the second is about the convergence from $M_n(\theta)$ to $M(\theta)$, which needs to be uniform.

Theorem 2.16. *Let* $\widehat{\theta}_n = \underset{\theta \in \Theta}{argmin}\, M_n(\theta)$ *and* $\theta_0 = \underset{\theta \in \Theta}{argmin}\, M(\theta)$. *For* $\Theta \subseteq \mathbb{R}^d$, *let* d *a distance on* $\mathbb{R}^d$. *Assume that the following two properties hold:*

1. *For all* $\varepsilon > 0$, $\sup_{d(\theta,\theta_0) \geq \varepsilon} M(\theta) > M(\theta_0)$.

2. $\sup_{\theta \in \Theta} |M_n(\theta) - M(\theta)| \xrightarrow{P} 0$.

Then $\widehat{\theta}_n \xrightarrow{P} \theta_0$.

Proof. From assumption 1., for every $\varepsilon > 0$, there exists η_ε such $d(\theta,\theta_0) \geq \varepsilon$ implies that $M(\theta) \geq M(\theta_0) + \eta_\varepsilon$. One can write

$$
\begin{aligned}
\mathbb{P}\left(d(\widehat{\theta}_n,\theta_0) \geq \varepsilon\right) &\leq \mathbb{P}\left(M(\widehat{\theta}_n) \geq M(\theta_0) + \eta_\varepsilon\right) \\
&= \mathbb{P}\left(\eta_\varepsilon \leq M(\widehat{\theta}_n) - M(\theta_0)\right) \\
&= \mathbb{P}\left(\eta_\varepsilon \leq M(\widehat{\theta}_n) - M_n(\widehat{\theta}_n) + M_n(\widehat{\theta}_n) - M_n(\theta_0) + M_n(\theta_0) - M(\theta_0)\right)
\end{aligned}
$$

As $\widehat{\theta}_n$ is a minimizer of M_n, we have $M_n(\widehat{\theta}_n) - M_n(\theta_0) \leq 0$ and

$$
\begin{aligned}
\mathbb{P}\left(d(\widehat{\theta}_n,\theta_0) \geq \varepsilon\right) &\leq \mathbb{P}\left(\eta_\varepsilon \leq M(\widehat{\theta}_n) - M_n(\widehat{\theta}_n) + M_n(\theta_0) - M(\theta_0)\right) \\
&\leq \mathbb{P}\left(\eta_\varepsilon \leq 2\sup_{\theta \in \Theta} |M_n(\theta) - M(\theta)|\right) \\
&= \mathbb{P}\left(\sup_{\theta \in \Theta} |M_n(\theta) - M(\theta)| \geq \frac{\eta_\varepsilon}{2}\right),
\end{aligned}
$$

and the right-hand side tends to zero by assumption 2., which concludes the proof.

$\square$

Remark 2.17. *Consistency also holds if* $\widehat{\theta}_n$ *is not an exact minimizer of* $M_n(\theta)$ *(which can be hard to compute in some practical cases), as long as its approximation error converges to zero (in probability). A sufficient condition to obtain consistency for an approximate minimizer is to further assume that*

$$
M_n(\widehat{\theta}_n) \leq M_n(\theta_0) + E_n
$$

for some random variable $E_n \xrightarrow{P} 0$.

Application to the MLE estimator   Using Theorem 2.16, we can propose some sufficient condition for the MLE to be a consistent estimator of $\theta_\star$ when $(X_1, \ldots, X_n) \sim P_{\theta_\star}$.

Theorem 2.18 (Consistency of the MLE)**.** *Assume that the model* $\mathcal{M} = \{f_\theta, \theta \in \Theta\}$ *satisfies the following properties:*

1. $\mathcal{M}$ *is identifiable, i.e.,* $f_\theta = f_{\theta'}$ *implies* $\theta = \theta'$ *for all* $(\theta, \theta') \in \Theta$.

2. Θ *is compact and for all* $x \in \mathcal{X}$, $\theta \mapsto f_\theta(x)$ *is continuous.*

3. *For all* $\theta \in \Theta$, $\mathbb{E}_\theta \left[\sup_{\theta' \in \Theta} |\log f_{\theta'}(X_1)| \right] < \infty$.

Then for all $\theta_\star \in \Theta$, *the MLE estimator built from a* n *sample* $X_1, \ldots, X_n \sim f_{\theta_\star}$ *satisfies* $\widehat{\theta}_n \xrightarrow{P} \theta_\star$ *(where the convergence is under the model* $\mathbb{P}_{\theta_\star}$).

For theses assumption to be satisfied in simple models such as Bernoulli and Gaussian, we would need to restrict the set of possible values for the means (to $[p_0, 1 - p_0]$ for $p_0 > 0$ in the Bernoulli case, or to some bounded interval $[a, b]$ in the Gaussian case). But in these two cases, the consistency of the MLE (which coincides with the empirical means) can already easily be established directly using the law of large number. Still a result such as Theorem 2.18 provide some generic guarantees for the MLE in potentially more complex models, under some restriction on the parameter space.

2.4.3 Asymptotic normality of the MLE estimator

Under stronger assumptions, it is also possible to further exhibit the limiting distribution of the MLE estimator. We start by presenting the result and a sketch of proof for the estimation of a one-dimensional parameter $\theta \in \mathbb{R}$. Given a n sample $X_1, \ldots, X_n \sim P_\theta$, we recall that $I(\theta)$ denotes the Fisher information obtained from one sample X_1.

Theorem 2.19. *Let* $\widehat{\theta}_n$ *be the MLE of a parameter* $\theta \in \mathbb{R}$ *computed on a* n *sample* $X_1, \ldots, X_n \sim P_\theta$. *If* $\widehat{\theta}_n$ *is consistent and if the model is regular (according to Definition 1.21) then if the Fisher information satisfies* $I(\theta) > 0$, $\sqrt{n}(\widehat{\theta}_n - \theta)$ *converges in distribution under* $\mathbb{P}_\theta$ *towards a Gaussian distribution:*

$$\sqrt{n}(\widehat{\theta}_n - \theta) \rightsquigarrow \mathcal{N}\left(0, \frac{1}{I(\theta)}\right) .$$

Proof. Let $\ell(\theta) = \log f_\theta(X_1, \ldots, X_n)$ be a simplified notation for the log-likelihood, and let $\ell'(\theta)$ be its derivative (in θ). As a minimizer of ℓ, the MLE estimator satisfies $\ell'(\widehat{\theta}_n) = 0$. Using a Taylor expansion of ℓ' in the true parameter θ, we can write

$$0 = \ell'(\widehat{\theta}_n) = \ell'(\theta) + (\widehat{\theta}_n - \theta)\ell''(\widetilde{\theta}_n)$$

for some $\widetilde{\theta}_n \in (\theta, \widehat{\theta}_n)$ (or $(\widehat{\theta}_n, \theta)$). Hence, one can write

$$\widehat{\theta}_n - \theta = -\frac{\ell'(\theta)}{\ell''(\widetilde{\theta}_n)}$$

$$\sqrt{n}\left(\widehat{\theta}_n - \theta\right) = \frac{\frac{1}{\sqrt{n}}\ell'(\theta)}{-\frac{1}{n}\ell''(\widetilde{\theta}_n)} \tag{2.2}$$

The numerator in (2.2) can be written

$$\frac{1}{\sqrt{n}}\ell'(\theta) = \sqrt{n}\left(\frac{1}{n}\sum_{i=1}^n s(X_i, \theta)\right) .$$

Using that under a regular model $\mathbb{E}_\theta\left[s(X_1; \theta)\right] = 0$ and $\text{Var}_\theta\left[s(X_1; \theta)\right] = I(\theta)$, one gets using the Central Limit Theorem that

$$\frac{1}{\sqrt{n}}\ell'(\theta) \rightsquigarrow \mathcal{N}\left(0, I(\theta)\right)$$

From the consistency of $\widehat{\theta}_n$ we know that $\widehat{\theta}_n \xrightarrow{P} \theta$, from which we deduce that $\widetilde{\theta}_n \xrightarrow{P} \theta$. The denominator of (2.2) can be written (this is the part where the proof becomes approximately correct)

$$-\frac{1}{n}\ell''(\widetilde{\theta}_n) = \frac{1}{n}\sum_{i=1}^{n} -\left(\frac{\partial^2 \log f_\theta(X_i)}{\partial^2 \theta}\right)_{\widetilde{\theta}_n} \simeq \frac{1}{n}\sum_{i=1}^{n} -\left(\frac{\partial^2 \log f_\theta(X_i)}{\partial^2 \theta}\right)_\theta$$

By the law of large numbers, under the model $\mathbb{P}_\theta$, this empirical average converges in probability to $\mathbb{E}_\theta\left[-\frac{\partial^2 \log f_\theta(X_1)}{\partial^2 \theta}\right]$ which is equal to the Fisher information $I(\theta)$ (using Lemma 1.26).

As $I(\theta) \neq 0$, we can use Slutsky's lemma to get that

$$\sqrt{n}\left(\widehat{\theta}_n - \theta\right) \sim \frac{1}{I(\theta)}\mathcal{N}\left(0, I(\theta)\right) = \mathcal{N}\left(0, \frac{1}{I(\theta)}\right).$$

□

A useful consequence of Theorem 2.16 is that it allows us to build asymptotic confidence regions around the MLE estimator, by replacing the (unknown) quantity $I(\theta)$ by its empirical version $I(\widehat{\theta}_n)$.

Corollary 2.20. *Under the assumptions of Theorem 2.19, if $\theta \mapsto I(\theta)$ is continuous in θ then under the model $\mathbb{P}_\theta$,*

$$\sqrt{nI(\widehat{\theta}_n)}\left(\widehat{\theta}_n - \theta\right) \rightsquigarrow \mathcal{N}(0,1).$$

Proof. Using the continuous mapping lemma and the consistency of $\widehat{\theta}_n$ yields that, under $\mathbb{P}_\theta$, $I(\widehat{\theta}_n) \xrightarrow{P} I(\theta)$. From Theorem 2.19 we have that under $\mathbb{P}_\theta$, $\sqrt{nI(\theta)}\left(\widehat{\theta}_n - \theta\right) \rightsquigarrow \mathcal{N}(0,1)$. Using Slutsky's lemma,

$$\sqrt{nI(\widehat{\theta}_n)}\left(\widehat{\theta}_n - \theta\right) = \sqrt{\frac{I(\widehat{\theta}_n)}{I(\theta)}} \times \sqrt{nI(\theta)}\left(\widehat{\theta}_n - \theta\right)$$

converges in distribution to $\mathcal{N}(0,1)$.

□

Hence, if our model is regular enough, we can build asymptotic confidence intervals of level $1 - \alpha$ around the MLE estimator (and resulting tests, see the next chapter) that are of the form

$$\left[\widehat{\theta}_n - \sqrt{\frac{1}{nI(\widehat{\theta}_n)}}q_{\alpha/2}; \widehat{\theta}_n - \sqrt{\frac{1}{nI(\widehat{\theta}_n)}}q_{\alpha/2}\right]$$

where q_α is such that $\mathbb{P}_{Z \sim \mathcal{N}(0,1)}\left(Z \leq q_\alpha\right) = 1 - \alpha$. We provide an example below.

Example 2.21. *Consider the Poisson model $X_1, \ldots, X_n \overset{iid}{\sim} \mathcal{P}(\lambda)$ where we recall that*

$$f_\lambda(k) = \frac{\lambda^k}{k!}e^{-\lambda}$$

for all $k \in \mathbb{N}$. The MLE is given by $\widehat{\lambda}_n = \frac{1}{n}\sum_{i=1}^{n} X_i$ and the Fisher information (of one sample) satisfies

$$I(\lambda) = \mathbb{E}_\lambda\left[-\frac{\partial^2}{\partial^2 \lambda}\log f_\lambda(X_1)\right]$$

We have

$$
\begin{aligned}
\frac{\partial \log f_\lambda(X_1)}{\partial \lambda} &= \frac{X_1}{\lambda} - 1 \\
\frac{\partial \log f_\lambda(X_1)}{\partial \lambda} &= -\frac{X_1}{\lambda^2}
\end{aligned}
$$

hence $I(\lambda) = \mathbb{E}_\lambda \left[\frac{X_1}{\lambda^2} \right] = \frac{1}{\lambda}$ *as the mean of a Poisson distribution with parameter* λ *is* λ. *Applying Corollary 2.20 yields that, under* $\mathbb{P}_\lambda$,

$$
\sqrt{\frac{n}{\widehat{\lambda}_n}} \left(\widehat{\lambda_n} - \lambda \right) \rightsquigarrow \mathcal{N}(0,1) .
$$

Now, let's use this information to build an asymptotic confidence interval on λ. *We have that*

$$
\begin{aligned}
\mathbb{P}_\lambda \left(-q_{\alpha/2} \le \sqrt{\frac{n}{\widehat{\lambda}_n}} \left(\widehat{\lambda_n} - \lambda \right) \le q_{\alpha/2} \right) \xrightarrow[n\to\infty]{} \ & \mathbb{P}_{Z\sim\mathcal{N}(0,1)} \left(-q_{\alpha/2} \le Z \le q_{\alpha/2} \right) \\
&= \ \mathbb{P}(Z \le q_{\alpha/2}) - \mathbb{P}\left(Z \le -q_{\alpha/2} \right) \\
&= \ \mathbb{P}(Z \le q_{\alpha/2}) - \mathbb{P}\left(Z > q_{\alpha/2} \right) \\
&= \ 1 - \frac{\alpha}{2} - \frac{\alpha}{2} \\
&= \ 1 - \alpha
\end{aligned}
$$

Putting λ *in the center of the interval, we have*

$$
\mathbb{P}_\lambda \left(\widehat{\lambda}_n - \sqrt{\frac{\widehat{\lambda}_n}{n}} q_{\alpha/2} \le \lambda \le \widehat{\lambda}_n + \sqrt{\frac{\widehat{\lambda}_n}{n}} q_{\alpha/2} \right) \xrightarrow[n\to\infty]{} 1 - \alpha
$$

which provides an asymptotic confidence interval of level $1 - \alpha$.

Extensions of Theorem 2.19 First, this result is also true for the estimation of a multi-dimensional parameter $\theta \in \mathbb{R}^d$ using the MLE. Under similar assumptions, we obtain that under $\mathbb{P}_\theta$,

$$
\sqrt{n}(\widehat{\theta}_n - \theta) \rightsquigarrow \mathcal{N}\left(0, I(\theta)^{-1} \right)
$$

but this time the Fisher information $I(\theta)$ is a $d \times d$ matrix, assumed to be invertible (see its definition in Section 1.4.2).

Then, while we presented the consistency results for the general family of M-estimators, we sticked to the MLE case for the asymptotic normality result. A counterpart of Theorem 2.19 also exists for M-estimators, see e.g. the book [Van der Vaart, 1998].

2.5 Asymptotic efficiency

In light of the Cramer-Rao lower bound given in Chapter 1, any estimator of a parameter $g(\theta) \in \mathbb{R}$ whose limit distribution satisfies

$$
\widehat{g}_n \approx \mathcal{N} \left(g(\theta), \frac{(g'(\theta))^2}{nI(\theta)} \right)
$$

is called *asymptotically efficient*. The reason is that, asymptotically, is is unbiased with a variance that is the minimal variance prescribed by the Cramer-Rao lower bound.

For estimating the parameter θ, (under appropriate regularity conditions) the MLE is an example of asymptotically efficient estimator, as we just saw that it satisfies

$$\widehat{g}_n \approx \mathcal{N}\left(\theta, \frac{1}{nI(\theta)}\right) \; .$$

However, we can find examples of MLE that are not efficient. Take for instance the MLE estimator of the variance of the Gaussian distribution, which is biased. We shall see other examples in exercises.

Chapter 3

Likelihood Ratio based Testing

After some reminder about the general formalism of (parameteric) statistical test, we present an important family of tests: likelihood ratio tests.

3.1 Statistical tests

Given a parametric model $X \sim P_\theta$, where $\theta \in \Theta$, a statistical test is an answer about some question about the unknown parameter θ of the form: does θ belong to a certain subset $\Theta_0 \subset \Theta$? Given two disjoint subsets Θ_0 and Θ_1 (that do not have to form a partition of Θ), a testing problem is characterized by two *hypotheses*:

$$\mathcal{H}_0 : (\theta \in \Theta_0) \quad \text{against} \quad \mathcal{H}_1 : (\theta \in \Theta_1).$$

When Θ_i is reduced to a singleton, i.e. $\mathcal{H}_i = (\theta = \theta_0)$, the hypothesis to test is called simple, otherwise it is called composite. $\mathcal{H}_0$ is called the null hypothesis and $\mathcal{H}_1$ the alternative hypothese. As we shall see, they play different roles.

A test can be formalized as a decision function $D : \mathcal{X} \to \{0, 1\}$, where $(D(X) = 1)$ means that given the observed data $X \in \mathcal{X}$ we tend to *reject* $\mathcal{H}_0$ and claim that $\mathcal{H}_1$ holds instead and $(D(X) = 0)$ means that we *do not reject* $\mathcal{H}_0$. Hence $(D(X) = 0)$ is not to be interpreted as a decision that $\mathcal{H}_0$ is true, rather that the observation X is still compatible with $\mathcal{H}_0$ being true.

The statistical guarantees that we can offer for a test are expressed in terms of type I error and type II error (or power).

- The type I error, denoted by α is often the primary concern of the statistician. It is defined for $\theta \in \Theta_0$ as

$$\alpha(\theta) = \mathbb{P}_\theta (D(X) = 1) \ .$$

  A small type I error means that the probability to reject $\mathcal{H}_0$ when it is actually true is small.

- The type II error at a given alternative $\theta \in \Theta_1$ is denoted by

$$\beta(\theta) = \mathbb{P}_\theta (D(X) = 0)$$

  and measure the probability to not detect that $\mathcal{H}_1$ in this alternative. Conversely, the *power* (in the alternative $\theta \in \mathcal{H}_1$) is the probability to detect that $\mathcal{H}_1$ holds, i.e. $1 - \beta(\theta)$.

In a statistical testing problem, we usually require to control the type I error for any possible $\theta \in \Theta_0$ and define the level of significance of the test to be $\alpha = \sup_{\theta \in \Theta_0} \alpha(\theta)$. Then, subject to this type I error control, we seek to provide guarantees on the type II error, or on the power, at least for θ in some part of the alternative.

In the (common) particular case where $X = (X_1, \ldots, X_n) \overset{iid}{\sim} P_\theta$ is a n samples, we can also consider the asymptotic quality of a test, when the sample size n goes large. In that case, indexing the test by its sample size, that is writing D_n instead of D, we define an asymptotic test of level α to be such that, for all $\theta \in \Theta_0$,

$$\lim_{n \to \infty} \mathbb{P}_\theta\left(D_n(X_1, \ldots, X_n) = 1\right) \le \alpha.$$

Similarly, one can look at the asymptotic type II error or power of a test.

Example and methodology   Tests are often based on some test statistic, denoted by $T = t(X) \in \mathbb{R}$ (T_n in the n-sample case) and reject the null hypothesis when T belongs to some rejection region $\mathcal{R} \subseteq \mathbb{R}$ (which typically depends on the desired significance level α we seek). We write

$$D(X) = \mathbb{1}\left(T \in \mathcal{R}\right) .$$

The test statistic is chosen so that its distribution (or its asymptotic distribution) under the null hypothesis (that is, for all $\theta \in \Theta_0$) is known. We use this knowledge to find a rejection region $\mathcal{R}_\alpha$ such that

$$\forall \theta \in \Theta_0, \quad \mathbb{P}_\theta\left(T \in \mathcal{R}_\alpha\right) \le \alpha$$

and our test at level α is $D_\alpha(X) = \mathbb{1}\left(T \in \mathcal{R}_\alpha\right)$. For power computation, it can be useful to also know the distribution of the test statistic under the alternative. In the n-sample setting, the power will typically increase with the sample size, and it is common to select the sample size so as to guarantee a minimal power in some part of the alternative.

Example 3.1. *We illustrate the testing methodology on a simple example: $X_1, \ldots, X_n$ is a n sample from a $\mathcal{N}(\mu, 1)$ distribution and we are testing the hypothesis*

$$\mathcal{H}_0 : (\mu = \mu_0) \quad against \quad \mathcal{H}_1 : (\mu = \mu_1)$$

for two distinct values $\mu_0 < \mu_1$.

The test is about the parameter μ so it is quite natural to base our test on an estimator of this parameter: we take the maximum likelihood estimator, which is the empirical mean $\widehat{\mu}_n = \frac{1}{n} \sum_{i=1}^n X_i$. In order to calibrate the test, we should know the distribution (or asymptotic distribution) of this estimator (or some transformation of it) under $\mathcal{H}_0$. In this simple Gaussian setting, $\widehat{\mu}_n$ also follows a Gaussian distribution: $\widehat{\mu}_n \sim \mathcal{N}\left(\mu_0, \frac{1}{n}\right)$. A natural test statistic is

$$T_n = \sqrt{n}\left(\widehat{\mu}_n - \mu_0\right)$$

as under $\mathcal{H}_0$, $T_n \sim \mathcal{N}(0, 1)$. Note that we also know the distribution of T_n under $\mathcal{H}_1$. Indeed by writing $T_n = \sqrt{n}(\mu_1 - \mu_0) + \sqrt{n}\left(\widehat{\mu}_n - \mu_1\right)$ we have that under $\mathcal{H}_1$, $T_n \sim \mathcal{N}\left(\sqrt{n}(\mu_1 - \mu_0), 1\right)$.

Given these observations, we expect T_n to take values close to zero when $\mathcal{H}_0$ is true, and very large values when $\mathcal{H}_1$ is true. This motivates a simple rejection region $\mathcal{R} = [t, +\infty[$ for some threshold t, hence the resulting test is $D_n(X) = \mathbb{1}(T_n > t)$. Now we calibrate the threshold in order to guarantee the desired type I error:

$$\mathbb{P}_{\theta_0}\left(D_n(X) = 1\right) = \mathbb{P}_{\theta_0}(T_n > t) = \alpha$$

if we choose $t = q_\alpha$, the $1 - \alpha$ quantile of the standard normal distribution: $\mathbb{P}_{Z \sim \mathcal{N}(0,1)}(Z \leq q_\alpha) = 1 - \alpha$. The type II error is

$$\mathbb{P}_{\theta_1}(D_n(X) = 0) = \mathbb{P}_{\theta_1}(T_n \leq q_\alpha) = \Phi(-\sqrt{n}(\mu_1 - \mu_0) + q_\alpha) = 1 - \Phi(\sqrt{n}(\mu_1 - \mu_0) - q_\alpha)$$

where Φ is the standard Gaussian cdf.

In order to guarantee a desired type II error β (additionally to the type I error which is already guaranteed for any sample size n), we can choose the sample size n to satisfy

$$\sqrt{n}(\mu_1 - \mu_0) - q_\alpha = q_\beta, \quad i.e. \quad n = \frac{(q_\alpha + q_\beta)^2}{(\mu_1 - \mu_0)^2}.$$

As we see in this example, finding good estimator and their distribution (or asymptotic distribution) under $\mathcal{H}_0$ and $\mathcal{H}_1$ is a very important first step in order to choose a test statistic and calibrate a test.

3.2 Likelihood-Ratio Tests

Given two hypotheses

$$\mathcal{H}_0 : (\theta \in \Theta_0) \quad \text{against} \quad \mathcal{H}_1 : (\theta \in \Theta_1).$$

the idea of a likelihood ratio test is to compare the likelihood of the observation $L(X; \theta)$ for parameters θ that come from Θ_0 and from Θ_1.

A Likelihood-Ratio Test (LRT) reject $\mathcal{H}_0$ for large enough values of the test statistic

$$\Lambda(X) = \frac{\sup_{\theta \in \Theta_1} L(X; \theta)}{\sup_{\theta \in \Theta_0} L(X; \theta)}$$

or equivalently for large values of the log-likelihood ratio

$$\log \Lambda(X) = \log \left(\frac{\sup_{\theta \in \Theta_1} L(X; \theta)}{\sup_{\theta \in \Theta_0} L(X; \theta)} \right).$$

More precisely, a LRT is of the form $D(X) = \mathbb{1}(\Lambda(X) > t)$ for some threshold $t > 1$ or $D(X) = \mathbb{1}(\log \Lambda(X) > u)$ for some $u > 0$. The idea of the test is to reject $\mathcal{H}_0$ when parameters in Θ_1 are significantly more likely that parameters in Θ_0.

Remark 3.2. *LRT are often used in settings where $\Theta_1 = \Theta \backslash \Theta_0$ where Θ is the entire set of possible parameters. In that case, one of the suprema will be equal to $L(X, \widehat{\theta}_{MLE})$. In particular, the likelihood-ratio test will reject $\mathcal{H}_0$ is the MLE belong to Θ_1 and satisfies*

$$\frac{L(X; \widehat{\theta}_{MLE})}{\sup_{\theta \in \Theta_0} L(X; \theta)} \geq t$$

for some well chosen threshold t.

Computational aspects In full generality, the computation of the likelihood ratio can be challenging, as unlike in the MLE case, we are required to solve at least one *constrained optimization* problem. In particular, when the hypotheses are composites (i.e. when neither Θ_0 nor Θ_1 are reduced to one element) the LRT is sometimes called the Generalized Likelihood Ratio Test (GLRT), the LRT being reserved to the simpler setting of testing two simple hypotheses where 1) the computation of the statistic is straightforward and 2) we have strong optimality guarantees, as will be explained in the next section.

Calibration of a LRT To use a LR test in practise, we need to find a threshold t_α such that if we define

$$D(X) = \mathbb{1}\left(\frac{\sup_{\theta \in \Theta_1} L(X; \theta)}{\sup_{\theta \in \Theta_0} L(X; \theta)} > t_\alpha\right)$$

the test has a type I error smaller than α for all $\theta \in \Theta_0$. This can be done case by case using arguments that are specific to the distribution at hand, but in Section 3.5 we will present a general result for calibrating (asymptotically) a GLRT.

Example 3.3. *Let's go back to the simple setting of Example 3.1. Letting f_μ be the density of a $\mathcal{N}(\mu, 1)$ distribution, the log-likelihood ratio in this case is*

$$
\begin{aligned}
\log \frac{\prod_{i=1}^n f_{\mu_1}(X_i)}{\prod_{i=1}^n f_{\mu_0}(X_i)} &= \sum_{i=1}^n \left[-\frac{(X_i - \mu_1)^2}{2} + \frac{(X_i - \mu_0)^2}{2}\right] \\
&= \frac{1}{2}\sum_{i=1}^n (\mu_1 - \mu_0)(2X_i - \mu_0 - \mu_1) \\
&= (\mu_1 - \mu_0)\left[n\widehat{\mu}_n - n\frac{\mu_0 + \mu_1}{2}\right].
\end{aligned}
$$

Hence, we see that $\log \Lambda(X) > u$ is equivalent to

$$\widehat{\mu}_n > \frac{\mu_0 + \mu_1}{2} + \frac{u}{n(\mu_1 - \mu_0)}.$$

Hence the LRT is of the form $D_n(X) = \mathbb{1}(\widehat{\mu}_n > x)$ for some threshold x. We remark that we this test is exactly the one proposed in Example 3.1, where we chose

$$x = \mu_0 + \frac{q_\alpha}{\sqrt{n}}$$

to guarantee a type I error α.

3.3 The Neyman-Pearson lemma

The result of this section is a first motivation for the use of LR tests: in some simple settings, LRT can be better than other tests, according to the following definition.

Definition 3.4. *Let $\alpha \in [0, 1]$. A statistical test D is called Uniformly More Powerful at level α (denoted by $\mathrm{UMP}(\alpha)$) if*

1. *the test D is of level α, i.e., $\sum_{\theta \in \Theta_0} \mathbb{P}_\theta(D(X) = 1) = \alpha$*

2. *For all other test D' that is of level α, $\forall \theta \in \Theta_1$, $\mathbb{P}_\theta(D(X) = 1) \geq \mathbb{P}_\theta(D'(X) = 1)$.*

The Neyman-Pearson lemma shows that for testing two simple hypotheses

$$\mathcal{H}_0 : (\theta = \theta_0) \quad \text{against} \quad \mathcal{H}_1 : (\theta = \theta_1)$$

where θ_0 and θ_1 are two distinct points of Θ, likelihood ratio tests, that have the simple form

$$D_t(X) = \mathbb{1}\left(\frac{f_1(X)}{f_0(X)} > t\right) \tag{3.1}$$

can be $\mathrm{UMP}(\alpha)$. Such simple likelihood-ratio tests are sometimes called Neyman-Pearson tests, due to the following result.

Theorem 3.5 (Neyman-Pearson lemma). *For $\alpha \in (0,1)$, if there exists a treshold $t_\alpha > 1$ such that the likelihood ratio test D_{t_α} as defined in* (3.1) *satisfies $\mathbb{P}_{\theta_0}(D_{t_\alpha}(X) = 1) = \alpha$, then this test is UPP($\alpha$).*

Proof. To ease the notation we denote by $D(X) = \mathbb{1}\left(\frac{f_{\theta_1}(X)}{f_{\theta_0}(X)} > t_\alpha\right)$ a LR test satisfying $\mathbb{P}_{\theta_0}(D(X) = 1) = \alpha$ and we let D' we another test satisfying $\mathbb{P}_{\theta_0}(D'(X) = 1) \leq \alpha$. In particular, we have

$$\mathbb{P}_{\theta_0}(D(X) = 1) - \mathbb{P}_{\theta_0}(D'(X) = 1) \geq 0 \quad \Leftrightarrow \quad \mathbb{E}_{\theta_0}[D(X) - D'(X)] \geq 0 \,.$$

Our goal is to prove that

$$\mathbb{P}_{\theta_1}(D(X) = 1) - \mathbb{P}_{\theta_1}(D'(X) = 1) \geq 0 \quad \Leftrightarrow \quad \mathbb{E}_{\theta_1}[D(X) - D'(X)] \geq 0 \,.$$

The proof consists in relating the expectation of $D(X) - D'(X)$ under θ_1 to that under θ_0. To do so, we introduce the function

$$g(x) = (D(x) - D'(x))(f_{\theta_1}(x) - t_\alpha f_{\theta_0}(x))$$

and we prove that this function is always non-negative by considering four cases:

1. If $f_{\theta_0}(x) = f_{\theta_1}(x) = 0$, then $g(x) = 0$.

2. If $f_{\theta_0}(x) = 0$ and $f_{\theta_1}(x) > 0$, then the value of the likelihood ratio $\frac{f_1(x)}{f_0(x)}$ is infinite, and as the threshold t_α is finite, we have $D(x) = 1$, which leads to $D(x) - D'(x) \geq 0$ for any D' and $g(x) = (D(x) - D'(x))f_{\theta_1}(x) \geq 0$.

3. If $f_{\theta_0}(x) > 0$ and $f_{\theta_1}(x) - t_\alpha f_{\theta_0}(x) > 0$, by definition of D, $D(x) = 1$ hence $D(x) - D'(x) \geq 0$ for any D' and $g(x) \geq 0$.

4. If $f_{\theta_0}(x) > 0$ and $f_{\theta_1}(x) - t_\alpha f_{\theta_0}(x) < 0$, by definition of D, $D(x) = 0$ hence $D(x) - D'(x) \leq 0$ for any D' and $g(x) \geq 0$.

We deduce that

$$\mathbb{E}_{\theta_1}[D(X) - D'(X)] - t_\alpha \mathbb{E}_{\theta_0}[D(X) - D'(X)] = \int_{\mathcal{X}} g(x)d\nu(x) \geq 0$$

hence $\mathbb{E}_{\theta_1}[D(X) - D(X')] \geq t_\alpha \mathbb{E}_{\theta_0}[D(X) - D'(X)] \geq 0$, which concludes the proof.

$\square$

The statement of Theorem 3.5 suggests that for a given $\alpha \in (0,1)$, there does not always exist a LR test that has a type I error exactly equal to α. This has to be nuanced a bit.

For continuous distributions, that is when f_{θ_0} is a density with respect to the Lebesgue measure, there is actually no such issue: it is always possible to find t_α such that

$$\mathbb{P}_{\theta_0}(f_1(X) > t_\alpha f_0(X)) = \alpha.$$

Indeed, letting $\mathcal{E}(t) = (f_1(X) > t f_0(X))$, one can justify that $t \mapsto \mathbb{P}(\mathcal{E}(t))$ is continuous, non-increasing and satisfies $\mathbb{P}(\mathcal{E}(0)) = 1$ and $\lim_{t \to \infty} \mathbb{P}(\mathcal{E}(t)) = 0$. By the intermediate value theorem, for any $\alpha \in (0,1)$, there exists t_α such that $\mathbb{P}(\mathcal{E}(t_\alpha)) = \alpha$.

For discrete distribution however, it is not always possible to exactly match some level α with a LR test. We can looke at a simple example to understand what happens: if we collect Bernoulli samples $X_1, \ldots, X_n \overset{iid}{\sim} \mathcal{B}(\theta)$ and we want to test $\mathcal{H}_0 : (\theta = \theta_0)$ and $\mathcal{H}_1 : (\theta = \theta_1)$ for $\theta_0 < \theta_1$, we can prove that a LR test is of the form $D_n(X) = \mathbb{1}\,(T(X) > t)$ where $T(X) = \sum_{i=1}^{n} X_i$ follows a binomial distribution $\mathcal{B}(n, \mu_0)$ under $\mathcal{H}_0$. As $T(X)$ only takes integer values in $\{0, 1, \ldots, n\}$, there are only $n + 1$ possible values of $\mathbb{P}_{\theta_0}(D_n(X) = 1)$, hence not all $\alpha \in (0, 1)$ can be exactly attained. For example if $\theta_0 = 0.6$ and $n = 10$ we can show that $\mathbb{P}_{\theta_0}(T(X) > 7) \simeq 0.167$ while $\mathbb{P}_{\theta_0}(T(X) > 8) \simeq 0.046$, hence there is no threshold that provides a type I error exactly equal to $\alpha = 0.05$. In that case, it is common to choose the threshold giving the largest type I error that is smaller than α, that is $t = 8$ in this example.

Remark 3.6 (randomized test)**.** *For discrete distributions, it is actually possible to exactly match a type I error α by considering the broader class of randomized tests. A randomized test is a mapping $\widetilde{D} : \mathcal{X} \to [0, 1]$ and $\widetilde{D}(X) = \gamma$ with $\gamma \in (0, 1)$ leads to rejecting $\mathcal{H}_0$ with probability γ. More concretely, the actual decision D to reject $\mathcal{H}_0$ or not reject it from the randomized test $\widetilde{D}$ can be written*

$$D(X) = \mathbb{1}(U \leq \widetilde{D}(X))$$

where $U \sim \mathcal{U}([0, 1])$ is a uniform random variable that is independent from X.

A randomized Likelihood Ratio test can be defined by a threshold t and a parameter $\gamma \in [0, 1)$ as follows:

$$\begin{aligned}
\widetilde{D}_{t,\gamma}(X) &= 1 & \text{if} \quad \frac{L(X; \theta_1)}{L(X; \theta_0)} > t \\
\widetilde{D}_{t,\gamma}(X) &= \gamma & \text{if} \quad \frac{L(X; \theta_1)}{L(X; \theta_0)} = t \\
\widetilde{D}_{t,\gamma}(X) &= 0 & \text{if} \quad \frac{L(X; \theta_1)}{L(X; \theta_0)} < t
\end{aligned}$$

The type I error of this randomized test is

$$\mathbb{P}_{\theta_0}(D_{t,\gamma}(X) = 1) = \mathbb{P}_{\theta_0}\left(\frac{L(X; \theta_1)}{L(X; \theta_0)} > t\right) + \gamma \mathbb{P}_{\theta_0}\left(\frac{L(X; \theta_1)}{L(X; \theta_0)} = t\right).$$

Hence in the discrete setting by choosing $t_\alpha = \inf\{t : \mathbb{P}_{\theta_0}\left(\frac{L(X; \theta_1)}{L(X; \theta_0)} > t\right) \leq \alpha\}$ and

$$\gamma_\alpha = \frac{\alpha - \mathbb{P}_{\theta_0}\left(\frac{L(X; \theta_1)}{L(X; \theta_0)} > t_\alpha\right)}{\mathbb{P}_{\theta_0}\left(\frac{L(X; \theta_1)}{L(X; \theta_0)} = t_\alpha\right)}$$

we end up with a test that has exactly a type I error α (and will have a larger power than when setting $\gamma = 0$).

By considering the more general class of randomized LRT, we can have a stronger version of Theorem 3.5 saying that for any $\alpha \in (0, 1)$ there exists a (possibly randomized) LRT with level α and that any other test of level α has a smaller power.

3.4 Particular forms of the Neyman-Person test

We already mentionned two examples (Gaussian and Bernoulli distribution) for which the Neyman-Pearson test ends up having a simple form. First, we explain that this properties can be extended to exponential families. We recall that a family of distributions P_θ forms an exponential families if their densities with respect to some common reference measure ν is of the form

$$f_\theta(x) = h(x) \exp\left(T(x)a(\theta) - b(\theta)\right) .$$

Proposition 3.7. *In an exponential family of the above form with $a(\theta_0) < a(\theta_1)$, the LR test for testing*

$$\mathcal{H}_0 : (\theta = \theta_0) \quad against \quad \mathcal{H}_1 : (\theta = \theta_1)$$

based on an n sample $X_1, \ldots, X_n$ from f_θ takes the form $D_n(X) = \mathbb{1}\left(\sum_{i=1}^n T(X_i) > t\right)$ where T is the canonical statistic of the exponential family.

Proof. The likelihood ratio can be written as follows:

$$\begin{aligned}
\Lambda(X) &= \frac{\prod_{i=1}^n h(X_i) \exp(T(X_i)a(\theta_1) - b(\theta_1))}{\prod_{i=1}^n h(X_i) \exp(T(X_i)a(\theta_0) - b(\theta_0))} = \frac{\exp(\sum_{i=1}^n T(X_i)a(\theta_0) - nb(\theta_0))}{\sum_{i=1}^n T(X_i)a(\theta_1) - nb(\theta_1)} \\
&= \exp\left((a(\theta_1) - a(\theta_0))\sum_{i=1}^n T(X_i) - n(b(\theta_1) - b(\theta_0))\right)
\end{aligned}$$

Hence $\log \Lambda(X) > u$ is equivalent to

$$(a(\theta_1) - a(\theta_0))\sum_{i=1}^n T(X_i) - n(b(\theta_1) - b(\theta_0)) > u .$$

As $a(\theta_1) - a(\theta_0) > 0$ this is in turn equivalent to

$$\sum_{i=1}^n T(X_i) > \frac{u + n(b(\theta_1) - b(\theta_0))}{a(\theta_1) - a(\theta_0)} .$$

$\square$

This property extends to any family of distribution that possesses a sufficient statistic (see Section 1.4.3). Recall that the density of a n samples $X_1, \ldots, X_n$ under such a distribution can be expressed as follows:

$$f_\theta(x_1, \ldots, x_n) = g(x_1, \ldots, x_n)h(S(x_1, \ldots, x_n); \theta) .$$

In that case $\Lambda(X) = \frac{h(S(X);\theta_1)}{h(S(X);\theta_0)}$ and the LR is of the form

$$D_n(X) = \mathbb{1}\left(S(X_1, \ldots, X_n) \in \mathcal{R}\right)$$

where the rejct region is of the form $\mathcal{R} = \{s \in \mathcal{Y} : \frac{h(s;\theta_1)}{h(s;\theta_0)} > t\}$ for some $t > 0$ and $S : \mathcal{X}^n \to \mathcal{Y}$.

3.5   Testing composite hypotheses

Simple hypotheses are quite restrictive. As a first step towards generalization, a more useful test is to check whether the value of the parameter θ is larger than some reference value. This corresponds to a case in which $\mathcal{H}_0$ is a simple hypothesis, while $\mathcal{H}_1$ is composite.

$$\mathcal{H}_0 : (\theta = \theta_0) \quad \text{against} \quad \mathcal{H}_1 : (\theta > \mu_0)$$

Such a test could be useful to access the efficacy of a new treatment in a setting where the average efficacy of the standard of care (θ_0) is considered known.

In Example 3.1, we studied a LR test for testing

$$\mathcal{H}_0 : (\mu = \mu_0) \quad \text{against} \quad \mathcal{H}_1 : (\mu = \mu_1) \tag{3.2}$$

based on iid observations $X_1, \ldots, X_n$ from $\mathcal{N}(\theta, 1)$, when $\mu_0 < \mu_1$. This test is given by

$$D_n(X) = \mathbb{1}\left(\widehat{\mu}_n > \mu_0 + \frac{q_\alpha}{\sqrt{n}}\right)$$

and has type I error α. As this test does not depend on μ_1, it is also a valid test for

$$\mathcal{H}_0 : (\mu = \mu_0) \quad \text{against} \quad \mathcal{H}_1 : (\mu > \mu_0) \ . \tag{3.3}$$

The type I error is still α (as $\mathcal{H}_0$ is the same). Now consider any other possible test D'_n for (3.3). For any $\mu_1 > \mu_0$, the Neyman-Pearson lemma applied to the test (3.2) tells us that $\mathbb{P}_{\mu_1}(D_n(X) = 1) \geq \mathbb{P}_{\mu_1}(D'_n(X) = 1)$. It follows that the test D_n is UMP(α) for (3.3).

In this simple example, one could investigate what a (Generalized) Likelihood Ratio would be for the composite hypothesis testing problem (3.3). The log-likelihood ratio is given by

$$\log \Lambda(X) = \sup_{\mu : \mu > \mu_0} \log \frac{L(X_1, \ldots, X_n; \mu)}{L(X_1, \ldots, X_n, \mu_0)}$$

and the same computations as in Example 3.3 further yield

$$\log \Lambda(X) = \sup_{\mu : \mu > \mu_0} (\mu - \mu_0)\left(n\widehat{\mu}_n - n\frac{\mu_0 + \mu}{2}\right)$$

To compute this constrained maximization, one can consider two cases: either the MLE $\widehat{\mu}_n$ is larger than μ_0 and the supremum is attained for $\mu = \widehat{\mu}_n$, or $\widehat{\mu}_n \leq \mu_0$ in which case for all $\mu > \mu_0$, $\widehat{\mu}_n - \frac{\mu_0 + \mu}{2} \leq 0$ and the function to maximize is always negative, and maximal at $\mu = \mu_0$, where it is zero. We obtain

$$\log \Lambda(X) = \frac{n}{2}(\widehat{\mu}_n - \mu_0)^2 \mathbb{1}(\widehat{\mu}_n \geq \mu_0) \ .$$

And $\log \Lambda(X) > u$ is equivalent to

$$\widehat{\mu}_n > \mu_0 + \sqrt{\frac{2u}{n}} \ ,$$

which is (again) the same form as the previous test. This provides an example of composite hypothesis testing problem in which the LRT is optimal, in that when calibrated to get a type I error α, it is UMP(α).

3.5.1 Optimality properties

Based on the previous example, one could hope that the optimality of LRT generalize to any composite hypothesis testing problem. Sadly, it is not the case. However, there exists some tests that are UMP(α) for particular hypothesis testing problems, under some conditions on the likelihood (and in some cases, those may coincide with LRTs, but there is no general rule about that). We present an example below.

Definition 3.8. *A family of distribution $\{P_\theta, \theta \in \Theta\}$ is a monotonic density ratio family if there exists a statistic $T(x)$ such that*

$$\forall \theta < \theta', \quad \frac{f_{\theta'}(x)}{f_\theta(x)} = g(T(x))$$

for some non-decreasing function g (that may depend on θ and θ').

Theorem 3.9 (Lehmann's theorem). *Consider the composite hypothesis testing*

$$\mathcal{H}_0 : (\theta \leq \theta_0) \quad against \quad \mathcal{H}_1 : (\theta > \theta_1)$$

based on a n-sample $X_1, \ldots, X_n$ whose (joint) distribution P_θ belongs to a monotonic density ratio family with statistic $T(x)$. Then the test given by $D_n(X) = \mathbb{1}(T(X_1, \ldots, X_n) > t)$ is UMP(α) where $\alpha = \sup_{\theta \leq \theta_0} \mathbb{P}_\theta(D_n(X) = 1)$.

Example 3.10. *The computations of Example 3.3 shows that the Gaussian n-sample model (with a fixed variance) is a monotonic density ratio family, with the statistic $T(x_1, \ldots, x_n) = \frac{1}{n} \sum_{i=1}^n X_i$. The same holds for exponential families, if $\theta \mapsto a(\theta)$ is increasing (see the proof of Proposition 3.7).*

3.5.2 Asymptotic calibration of a Likelihood-Ratio Test

In this section, we discuss Likelihood Ratio Tests for a generic hypothesis testing problem of the form

$$\mathcal{H}_0 : (\theta \in \Theta_0) \quad \text{against} \quad \mathcal{H}_1 : (\theta \in \Theta \backslash \Theta_0) \tag{3.4}$$

(that is, we have $\Theta_1 = \Theta \backslash \Theta_0$). In this context, the likehood ratio as we defined it should be written

$$\Lambda(X) = \frac{\sup_{\theta \in \Theta \backslash \Theta_0} L(X; \theta)}{\sup_{\theta \in \Theta_0} L(X; \theta)} .$$

For when $\Theta_1 = \Theta \backslash \Theta_0$, it is more common to define

$$\widetilde{\Lambda}(X) = \frac{\sup_{\theta \in \Theta} L(X; \theta)}{\sup_{\theta \in \Theta_0} L(X; \theta)} .$$

First, from a computational perspective, the latter is preferable as it only features one constrained optimization problem. Then, it can be observed that for any threshold $t > 1$,

$$\Lambda(X) > t \quad \Leftrightarrow \quad \widetilde{\Lambda}(X) > t$$

so the decision of the associated test is the same. To justify this, we observe that $\Lambda(X) > 1$ if an only if the MLE estimator $\widehat{\theta} = \operatorname*{argmax}_{\theta \in \Theta} L(X; \theta)$ belongs to $\Theta \backslash \Theta_0$, and in that case $\Lambda(X) = \widetilde{\Lambda}(X)$.

Even if there are no general optimality properties (expressed with the UMP(α) property), we show below a general result providing a way to calibrate a Generalized Likelihood Ratio test, using asymptotic considerations.

Theorem 3.11 (Wilk's theorem). *Consider a n-samples $X_1, \ldots, X_n \sim \mathbb{P}_\theta$ coming from a parameteric model $\{\mathbb{P}_\theta, \theta \in \Theta\}$ with $\dim(\Theta) = p$. Let $\widetilde{\Lambda}$ be the likelihood-ratio associated to the test (3.4) where Θ_0 is a subspace of Θ with $\dim(\Theta_0) = q$. Then, for any $\theta \in \Theta_0$,*

$$2 \log \widetilde{\Lambda}(X_1, \ldots, X_n) \rightsquigarrow \chi^2(p-q)$$

Hence the test $D_n(X) = \mathbb{1}\left(2 \log \widetilde{\Lambda}(X_1, \ldots, X_n) > t_\alpha\right)$ with t_α equal to the $1 - \alpha$ quantile of the chi-square distribution with $p - q$ degrees of freedom is asymptotically of level α.

Bibliography

[Rivoirard and Stoltz, 2009] Rivoirard, V. and Stoltz, G. (2009). *Statistique en Action*. Vuibert.

[Van der Vaart, 1998] Van der Vaart, A. (1998). *Asymptotic Statistics*. Cambridge University Press.

[Wasserman, 2004] Wasserman, L. (2004). *All of Statistics, A Concise Course in Statistical Inference*. Springer.